

# Employee Absenteeism

*Gideon Anthony D'Souza*

Jan 2019

---

# Contents

---

<b>Contents</b>	<b>2</b>
<b>1 Introduction</b>	<b>4</b>
1.1 Problem Statement . . . . .	4
1.2 Data . . . . .	4
1.2.1 Data Overview: . . . . .	5
1.2.2 Data Description . . . . .	6
<b>2 Data Pre-Processing</b>	<b>8</b>
2.1 Data Preparation . . . . .	8
2.2 Missing Value Analysis . . . . .	9
2.3 Outlier Analysis . . . . .	13
2.4 Data Understanding . . . . .	15
<b>3 Feature Selection</b>	<b>21</b>
3.0.1 Correlation Analysis . . . . .	21
3.0.2 Analysis Of Variance (ANOVA) . . . . .	21
<b>4 Feature Engineering</b>	<b>24</b>
4.0.1 Feature Scaling: . . . . .	24
4.0.2 Creation of dummy variables . . . . .	24
<b>5 Modeling</b>	<b>27</b>
5.1 Model Selection . . . . .	27
5.2 Methodology: . . . . .	27
5.2.1 Hold-Out Method . . . . .	27
5.2.2 Hyper parameter Tunning . . . . .	28
5.2.3 K-Fold Cross Validation Technique . . . . .	28
5.3 Regression Model Building . . . . .	28
5.3.1 Backward Elemination: . . . . .	28
5.3.2 Principal component analysis: . . . . .	33
5.4 Classification Model Building . . . . .	36

<b>6 Conclusion</b>	<b>54</b>
6.1 Model Evaluation . . . . .	54
6.2 Model Selection . . . . .	54
6.3 What changes company should bring to reduce the number of absenteeism? . . . . .	55
6.4 How much losses every month can we project in 2011 if same trend of absenteeism continues? . . . . .	55
<b>Appendices</b>	<b>57</b>
<b>Appendices</b>	<b>58</b>
<b>A R code snapshots:</b>	<b>58</b>

## Chapter 1

# Introduction

## 1.1 Problem Statement

The objective of this Case is as follows.

- What changes company should bring to reduce the number of absenteeism?
- How much losses every month can we project in 2011 if same trend of absenteeism continues?

## 1.2 Data

In order to predict future absenteeism based on the given data, it is required to build a regressor model on the data set provided. Table 1.1, 1.2, Table 1.3 & 1.6 below is a sample of the data set that is using to predict employee absenteeism.

Table 1.1: Absenteeism (Columns(1-6))

ID	Reason for absence	Month of absence	Day of the week	Seasons	Transportation expense
11	26.0	7.0	3	1	289.0
36	0.0	7.0	3	1	118.0
3	23.0	7.0	4	1	179.0
7	7.0 7.0	5	1	1	279.0
11	23.0	7.0	5	1	289.0

Table 1.2: Absenteeism (Columns(7-11))

Distance from Residence to Work	Service time	Age	Work load Average/day	Hit target
36.0	13.0	33.0	239554.0	97.0
13.0	18.0	50.0	239554.0	97.0
51.0	18.0	38.0	239554.0	97.0
5.0	14.0	39.0	239554.0	97.0
36.0	13.0	33.0	239554.0	97.0

Table 1.3: Absenteeism (Columns(12-17))

<b>Disciplinary failure</b>	<b>Education</b>	<b>Son</b>	<b>Social drinker</b>	<b>Social smoker</b>	<b>Pet</b>
0.0	1.0	2.0	1.0	0.0	1.0
1.0	1.0	1.0	1.0	0.0	0.0
0.0	1.0	0.0	1.0	0.0	0.0
0.0	1.0	2.0	1.0	1.0	0.0
0.0	1.0	2.0	1.0	0.0	1.0

Table 1.4: Absenteeism (Columns(18-22))

<b>Weight</b>	<b>Education</b>	<b>Height</b>	<b>Body mass index</b>	<b>Absenteeism time in hours</b>
90.0	172.0	30.0	4.0	
98.0	178.0	31.0	0.0	
89.0	170.0	31.0	2.0	
68.0	168.0	24.0	4.0	
90.0	172.0	30.0	2.0	

### 1.2.1 Data Overview:

As shown in the Table 1.7 there are 21 features or predictor variables to predict the employee absenteeism which is the target or label. There is one target variables as shown in Table 1.7.

Table 1.5:

<b>S.No</b>	<b>Predictor</b>	<b>Target</b>
1	ID	Absenteeism time in hours
2	Reason for absence	
3	Month of absence	
4	Day of the week	
5	Seasons	
6	Transportation expense	
7	Distance from Residence to Work	
8	Service time	
9	Age	
10	Work load Average/day	
11	Hit target	
12	Disciplinary failure	
13	Education	
14	Son	
15	Social drinker	
16	Social smoker	
17	Pet	
18	Weight	
19	Education	
20	Height	
21	Body mass index	

## 1.2.2 Data Description

From the given in the metadata file, the feature and target variables names and descriptions are given in table 1.6

Table 1.6: Feature and target description. Detailed description for features with \*\*

S.No	Predictor/target variables	Description
1	ID	Employee ID
2	Reason for absence	Reason for absence**
3	Season	Season (1:spring, 2:summer, 3:fall, 4:winter)
4	Transportation expense	travel expense
5	Month of absence	Month (1 to 12)
6	Distance from Residence to Work	Distance in km
7	Day of the week	Monday - Friday (2-5)
8	Service time	Years employed in company.
9	Age	Age of employee
10	Work load Average/day	work load per day
11	Hit target	target achieved in percentage
12	Disciplinary failure	Disciplinary failure(0-no 1-yes)
13	Education	Education.**
14	Son	number of sons.
15	Pet	number of pets.
16	Social drinker	0 - no, 1 - yes.
17	Social smoker	0 - no, 1 - yes.
18	Weight	employee in kg.
19	Height	Height of employee.
20	Body mass index	employee BMI.
21	Absenteeism time in hours	Absenteeism in hours

### – Reason for absence:

- \* Certain infectious and parasitic diseases
- \* Neoplasms
- \* Diseases of the blood and blood-forming organs and certain disorders involving the immune mechanism
- \* Endocrine, nutritional and metabolic diseases
- \* Mental and behavioral disorders
- \* Diseases of the nervous system
- \* Diseases of the eye and adnexaVIII Diseases of the ear and mastoid process
- \* Diseases of the circulatory system
- \* Diseases of the respiratory system
- \* Diseases of the digestive system
- \* Diseases of the skin and subcutaneous tissue
- \* Diseases of the musculoskeletal system and connective tissue
- \* Diseases of the genitourinary system

- \* Pregnancy, childbirth and the puerperium
- \* Certain conditions originating in the perinatal period
- \* Congenital malformations, deformations and chromosomal abnormalities
- \* Symptoms, signs and abnormal clinical and laboratory findings, not elsewhere classified
- \* Injury, poisoning and certain other consequences of external causes
- \* External causes of morbidity and mortality
- \* Factors influencing health status and contact with health services.
- \* medical consultation
- \* blood donation
- \* laboratory examination
- \* unjustified absence
- \* physiotherapy
- \* dental consultation.

– **Education:**

high school (1), graduate (2), postgraduate (3), master and doctor (4).

From the information provided by the metadata. The continuous and categorical feature variables can clearly be separated. Table 1.7 shows the categorical and continuous variables provided in the data set.

Table 1.7: List of categorical and continuous features present in the data set.

S.No	Categorical features	Continuous features
1	Season	Transportation expense
2	ID	Distance from Residence to Work
3	Reason for absence	Service time
4	Month of absence	Age
5	Day of the week	Work load Average/day
6	Disciplinary failure	Hit target
7	Education	Weight
8	Son	Height
9	Pet	Body mass index
10	Social drinker	Absenteeism time in hours
11	Social smoker	

---

# Data Pre-Processing

---

## 2.1 Data Preparation

In order to create a model for the data it is first required to look into the data. First the data types of the feature and target variables are checked in order to see if it complies with the information given in the meta data file (data description file). In order to perform the aforementioned the variables are split according to their data types by the function listing ??.

Listing 2.1: Function to separate the variables names into their respective data types.

```
1 def dtype_separator(df, col_names):
2     '''
3     #####↵
4     This function will segregate the different data types present in ↵
5     the dataframe.
6     #####↵
7     Input -
8     * df - The data frame to be analysed.
9     *A list of the column names.
10    Output -
11    * obj - columns containing object data type.
12    * num - columns containing numerical data type (int64/float64).
13    * bool_d - columns containing bool data type.
14    * unknown = columns containing data type such as datetime64, ↵
15    timedelta[ns] and others.
16    #####↵
17    '''
18    obj = []
19    num = []
20    bool_d = []
21    unknown = []
22    for i in range(len(col_names)):
```



```

21         if df.iloc[:,i].dtype == 'O':
22             obj.append(col_names[i])
23         elif df.iloc[:,i].dtype == 'int64' or df.iloc[:,i].dtype == 'float64':
24             num.append(col_names[i])
25         elif df.iloc[:,i].dtype == 'bool':
26             bool_d.append(col_names[i])
27         else:
28             unknown.append(col_names[i])
29     return obj, num, bool_d, unknown

```

---

Function ?? outputs list of columns according to their data types. It was noted that the categorical features were incorrectly saved as numerical data type and hence they are to be converted back to categorical data type. It was also noted that the months feature ranges from 0-12, hence the value 0 is converted to NA as there is only 12 months in a year. Work load average is divided by 1000 as it was informed by the support team i a ticket raised by me.

## 2.2 Missing Value Analysis

The data set was checked for missing values and was found that Body mass index has the highest percentage of missing value of 4%. The missing value percentage for each feature is shown in figure 2.1. In order to deal with the missing values we impute them using KNN with a k value of 3. This would compute the k nearest values and replace the na with the respective k nearest value.

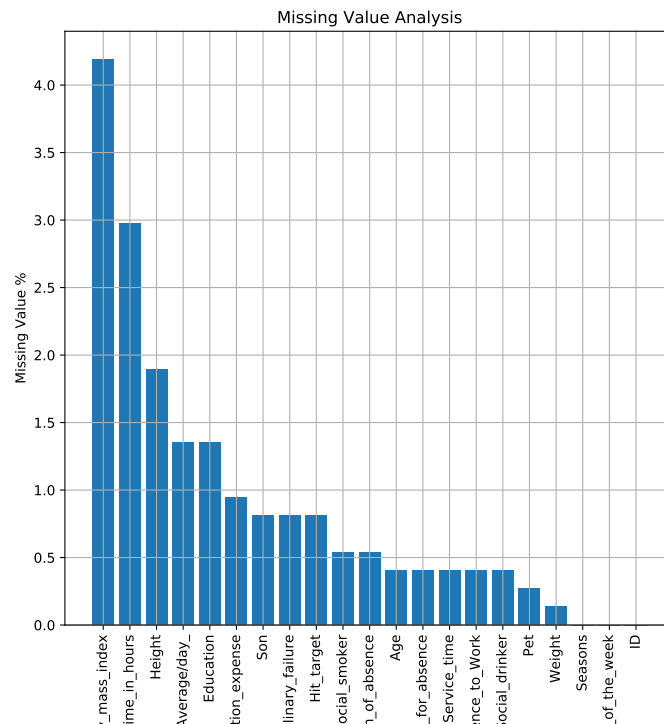


Figure 2.1: Missing value percentage for each feature.

After conducting the missing value analysis, the distributions of the continuous features are plotted. This is shown in figure 2.2, 2.3, 2.4, 2.5 & 2.6. Figure 2.7 is the distribution plot of the target variables. It can be noted that the data is not nominally distributed and after analysis the skewness it can be seen that there is a slight skew as shown in figure 2.8. We will conduct an outlier analysis and see if this affects the distributions.

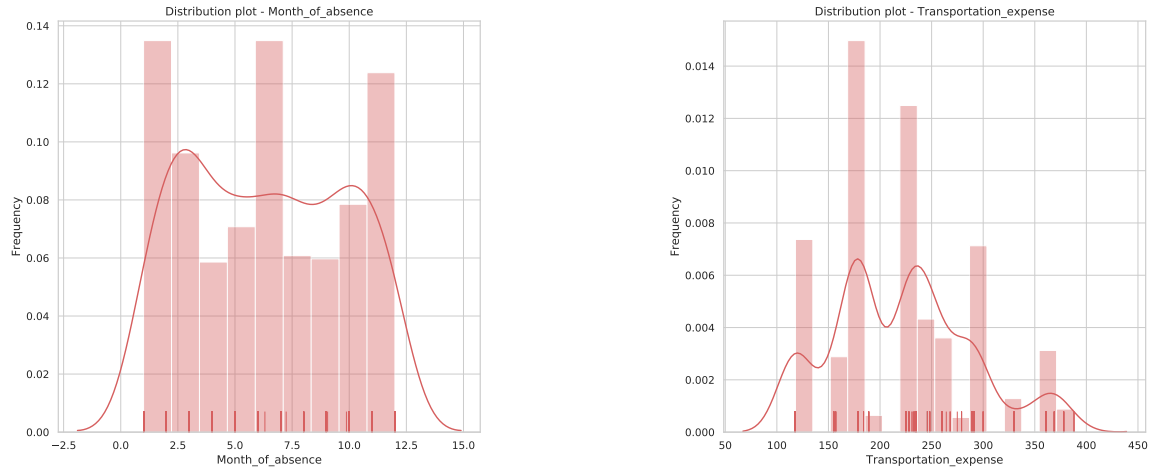


Figure 2.2: Distribution plots of features before outlier analysis.

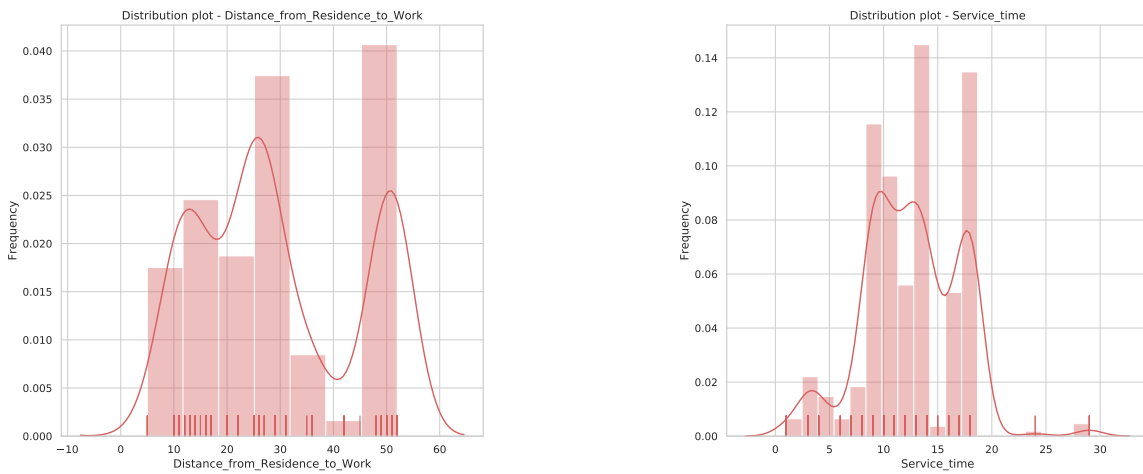


Figure 2.3: Distribution plots of features before outlier analysis.

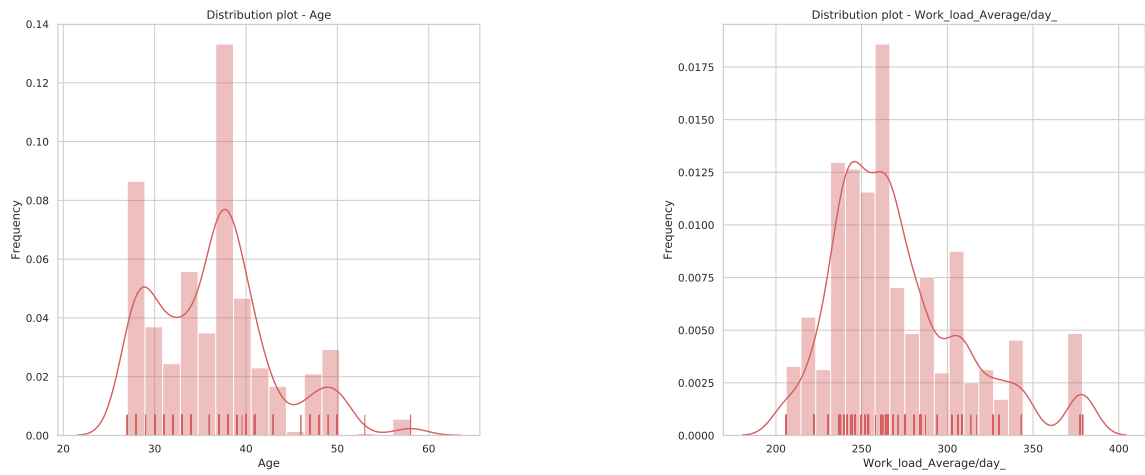


Figure 2.4: Distribution plots of features before outlier analysis.

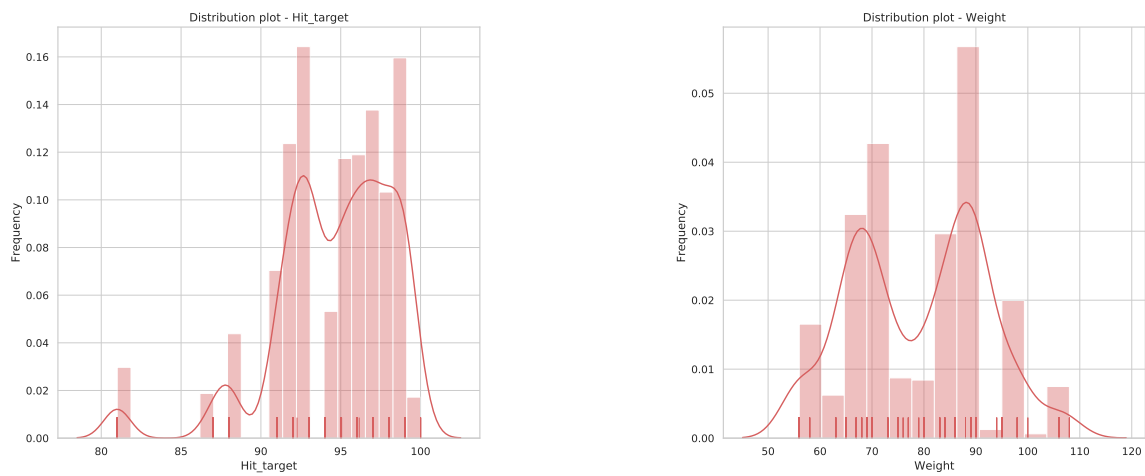


Figure 2.5: Distribution plots of features before outlier analysis.

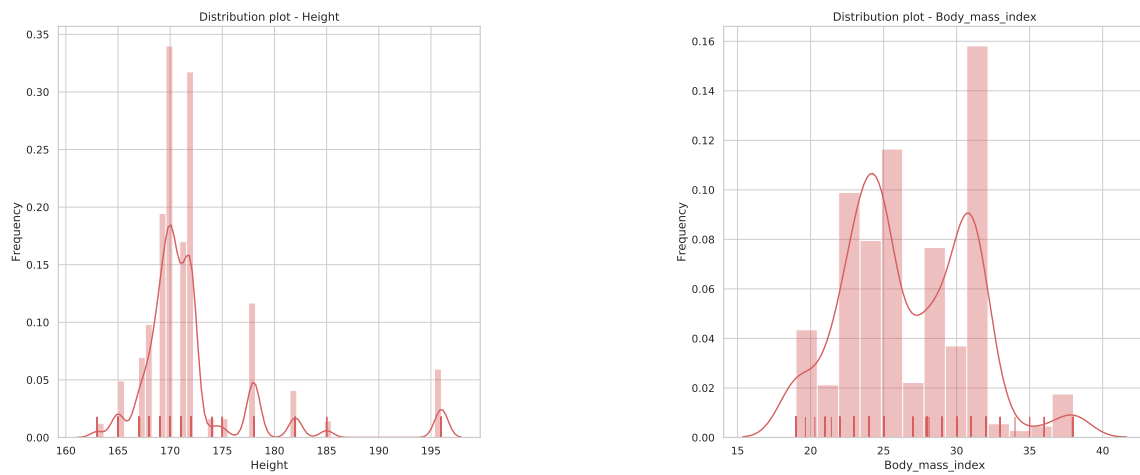


Figure 2.6: Distribution plots of features before outlier analysis.

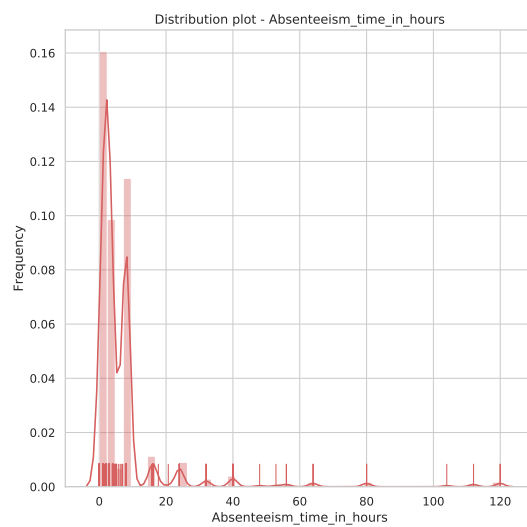


Figure 2.7: Distribution plots of target before outlier analysis.

```

*****
statistical properties of Month of absence:
DescribeResult(nobs=740, minmax=(1.0, 12.0), mean=6.354700578338823, variance=11.65477971915844, skewness=0.070265400731
89858, kurtosis=-1.2589112442736288)
*****
statistical properties of Transportation expense:
DescribeResult(nobs=740, minmax=(118.0, 388.0), mean=221.31365847071507, variance=4480.378835673238, skewness=0.39566180
701090514, kurtosis=-0.3221653522364045)
*****
statistical properties of Distance from Residence to Work:
DescribeResult(nobs=740, minmax=(5.0, 52.0), mean=29.631081080006933, variance=220.1302910973327, skewness=0.31144982841
84064, kurtosis=-1.2612709700406315)
*****
statistical properties of Service time:
DescribeResult(nobs=740, minmax=(1.0, 29.0), mean=12.554054054054054, variance=19.22711480086311, skewness=-0.0047099912
53832424, kurtosis=0.6704067197065524)
*****
statistical properties of Age :
DescribeResult(nobs=740, minmax=(27.0, 58.0), mean=36.44789531460408, variance=41.95812343690743, skewness=0.69731597063
86888, kurtosis=0.42388967427125657)
*****
statistical properties of Work_load Average/day :
DescribeResult(nobs=740, minmax=(205.917, 378.884), mean=271.1603506483232, variance=1507.4032516337088, skewness=0.9711
718700009127, kurtosis=0.6453444214427804)
*****
statistical properties of Hit target:
DescribeResult(nobs=740, minmax=(81.0, 100.0), mean=94.58582517619392, variance=14.28251421473298, skewness=-1.258105085
860628, kurtosis=2.39243588672608)
*****
statistical properties of Weight:
DescribeResult(nobs=740, minmax=(56.0, 108.0), mean=79.03513513491443, variance=165.9771129815394, skewness=0.0169668903
97851576, kurtosis=-0.9158612722683972)
*****
statistical properties of Height:
DescribeResult(nobs=740, minmax=(163.0, 196.0), mean=172.11497717710074, variance=36.42001797662386, skewness=2.56095929
61306387, kurtosis=7.260198747582171)
*****
statistical properties of Body_mass index:
DescribeResult(nobs=740, minmax=(19.0, 38.0), mean=26.6813915575566, variance=18.366170391434352, skewness=0.29994337932
45435, kurtosis=-0.32299568705714377)
*****
statistical properties of Absenteeism time in hours:
DescribeResult(nobs=740, minmax=(0.0, 120.0), mean=7.018145823809551, variance=179.92603279103142, skewness=5.6323240010
94293, kurtosis=37.50800276377008)
*****

```

Figure 2.8: Skwness test before outlier analysis.

## 2.3 Outlier Analysis

From figure 2.2, 2.3, 2.4, 2.5, 2.6, 2.7 & ?? it can be observed that most of the variables are skewed. This could be due to outliers present in the data set. In order to visualize outliers, the classic approach of using boxplot is used. Figure 2.9 & 2.10 Shows the outlier analysis using boxplot. It can be seen from figure 2.9 there that hight variables have outliers present in them and figure 2.10 shows the presence of outliers in the target variable. **Note: Outlier where also present in other features and their output is attached in the appendix.** In order to decide as to how to deal with the outliers in the data set, the percentage of outliers in the data set is first calculated by sample code in listing 2.2. the output of listing 2.2 shows that the **percentage of outliers present in the data set in 31 %**. In such a case it was decided to use knn imputation to compute the values of the outliers and replace them with it.

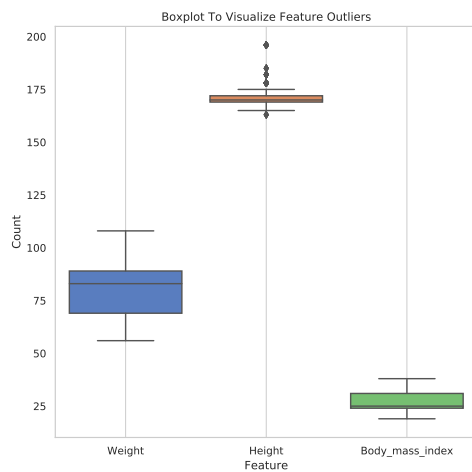


Figure 2.9: Outliers present in the feature variables.

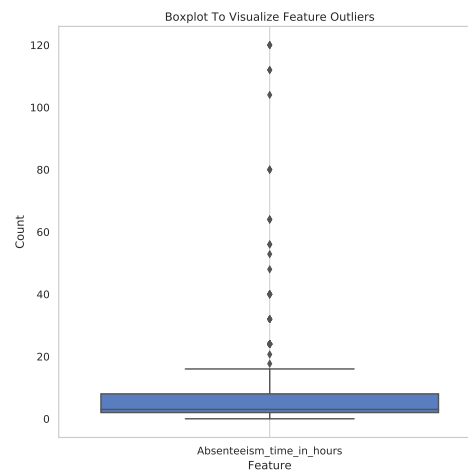


Figure 2.10: Outliers present in the target variables.

Listing 2.2: Function to calculate outlier percentage.

---

```

1 l_q = df.quantile(0.25)
2 u_q = df.quantile(0.75)
3 iqr = u_q - l_q
4 total_outliers = ((df<(l_q - 1.5*iqr))|(df>(u_q + 1.5*iqr))).sum()
5 outliers = total_outliers.sum()
6
7 r,c = df.shape
8 outlier_percentage = (outliers/r)*100
9 print("Outlier percentage :",round(outlier_percentage,3))

```

---

Now that the outliers are replaced with their respective knn results, the distributions of the continuous variables are compared. For the purpose of the report the distribution for absenteeism time in hours is analyzed. The rest of the distributions can be seen in the appendix section. From figure ?? it can be seen that after outlier analysis, the maximum value of the distribution is reduced from 120 to 16. It can also be noted that the shape of the distribution is changed. The distribution still does not resemble a normal distribution. The distribution seems more likely to be Poisson in nature.

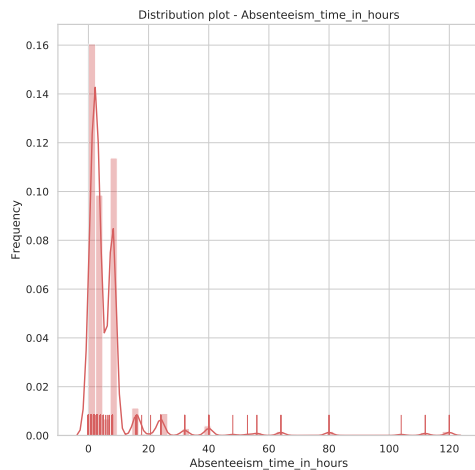


Figure 2.11: Distribution of the target variable before outlier analysis.

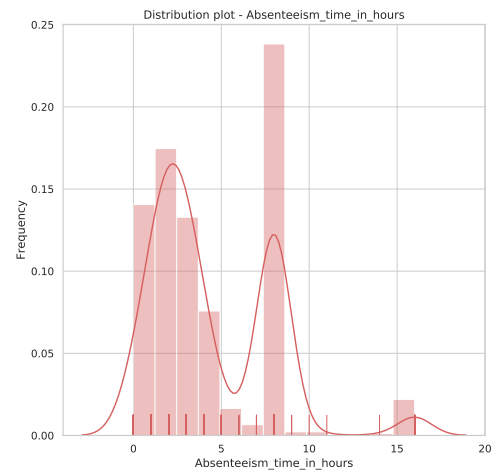


Figure 2.12: Distribution of the target variable before outlier analysis.

## 2.4 Data Understanding

In order to get further insight and understand of the data set. It required is to see as to how different features interact with each other and the target. First the amount of absenteeism counts verses employee id is analyzed. this is shown in figure 2.13. It can be seen that employee 9 and 28 has the maximum absenteeism count for between the time 0-5 hours and the reason is for medical consultation. it is also noted that employee id 17, 11 and 24 takes the maximum amount of unjustified absence and are mainly absent for the whole day that is 10 hours.

Next we shall analyses the education and absenteeism count along side the reason. From figure 2.4 it can be said that employees with only a high school degree tend to be absent the most and the most often used reason is medical consultation. Where as other employees with higher qualification are absent significantly less.

Now we check as to how absenteeism varies from pet owners and people who have children. From figure 2.4 By plotting number of children verses absentees it can be seen that employees with no children tend to be absent more often that the others. Next by plotting number of pets verses absentees it can be seen that employees with no pets tend to be absent more often that the others.

Next we analysis the month of absence with absenteeism count. From 2.18 we can see that the employee absentees is almost the same through the year.

Next we analysis the work load per day vs absenteeism count with respect to service time and age. From 2.19 we can see that the employee having medium amount of work load i.e.240 - 300 work load per day and who are below the age of 20 and having service time less than 8 years tend to be absent more often.

We now analysis the absenteeism vs month of absenteeism with respect to education. From 2.20 it is clearly seen that employees having only a high school tend to be absent more often than the other employees.

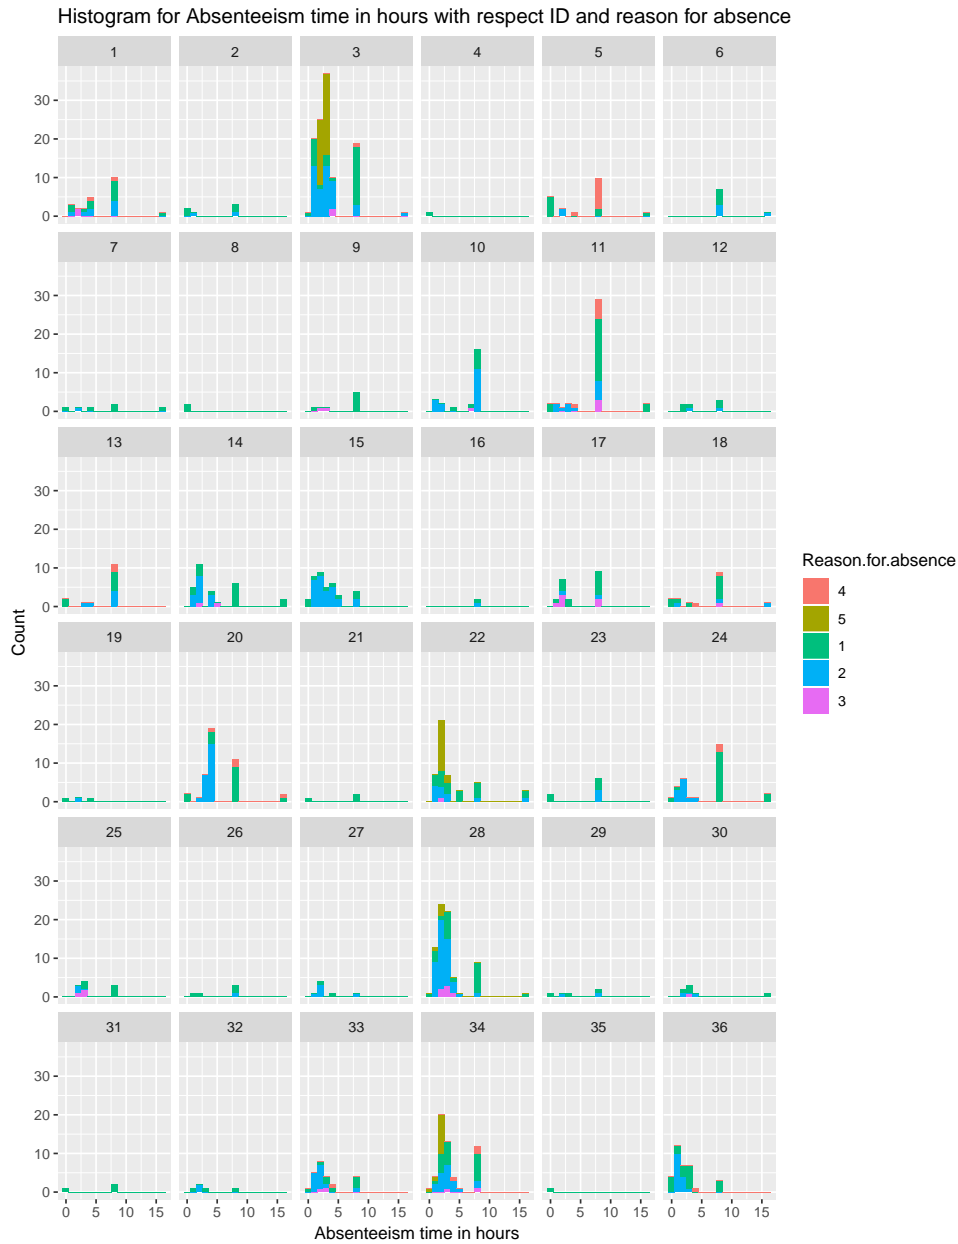


Figure 2.13: Analysis between employee id and absenteeism count. It can be seen that employee 9 and 28 has the maximum absenteeism count for between the time 0-5 hours and the reason is for medical consultation. it is also noted that employee id 17, 11 and 24 takes the maximum amount of unjustified absence and are mainly absent for the whole day that is 10 hours. The categories are as follows 1 - Code of Diseases, 2 - medical consultation, 3 - laboratory consultation, 4 - unjustified absence and 5 - physiotherapy



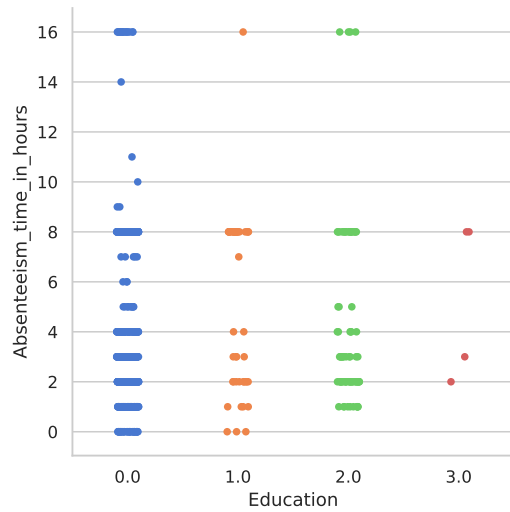


Figure 2.14: By plotting education qualification verses absenteeism it can be seen that employees with only a high school qualification tend to be absent more often that the others. Education (high school (1), graduate (2), postgraduate (3), master and doctor (4))

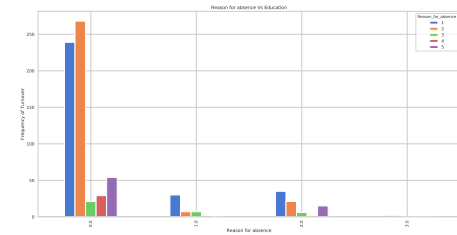


Figure 2.15: By plotting education qualification verses absentees with respect to the reason of absenteeism. it can be said that employees with only a high school degree tend to be absent the most and the most often used reason is medical consultation. Where as other employees with higher qualification are absent significantly less.

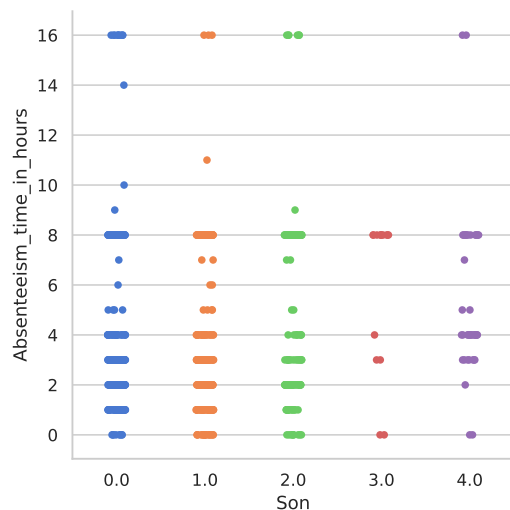


Figure 2.16: By plotting number of children verses absenteeism it can be seen that employees with no children tend to be absent more often that the others.

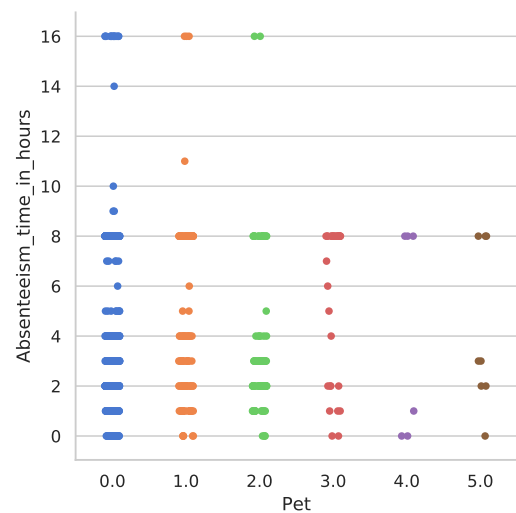


Figure 2.17: By plotting number of pets verses absenteeism it can be seen that employees with no pets tend to be absent more often that the others.

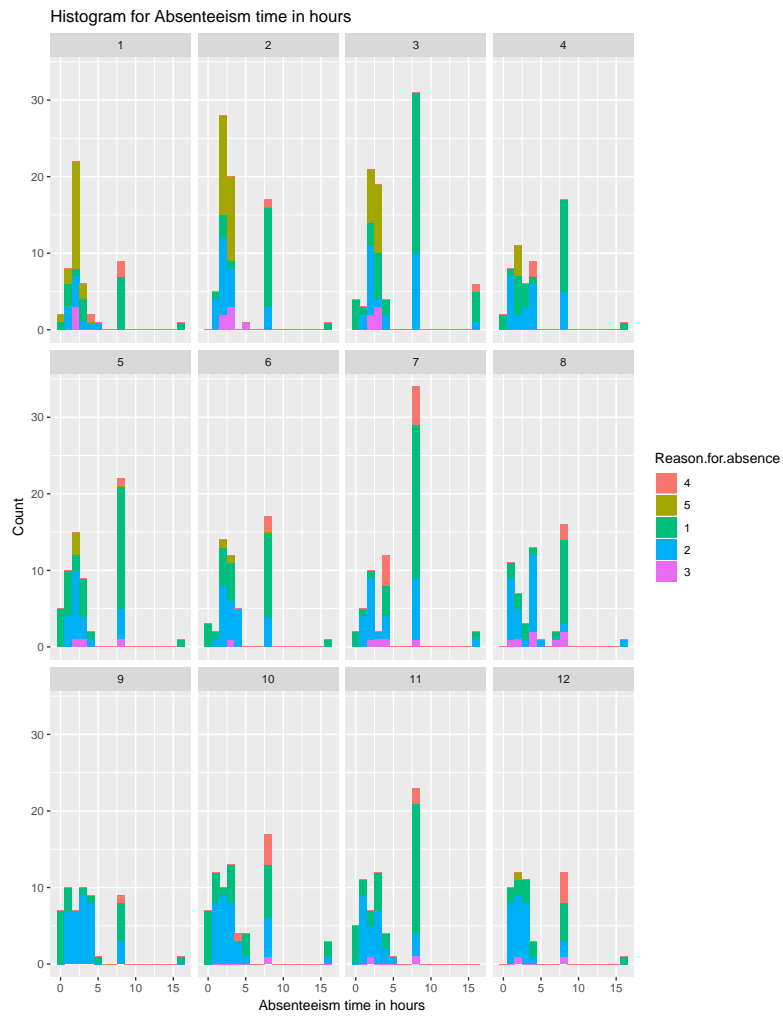


Figure 2.18: By plotting the month of absence verses reason. we can see that the employee absentees is almost the same through the year.

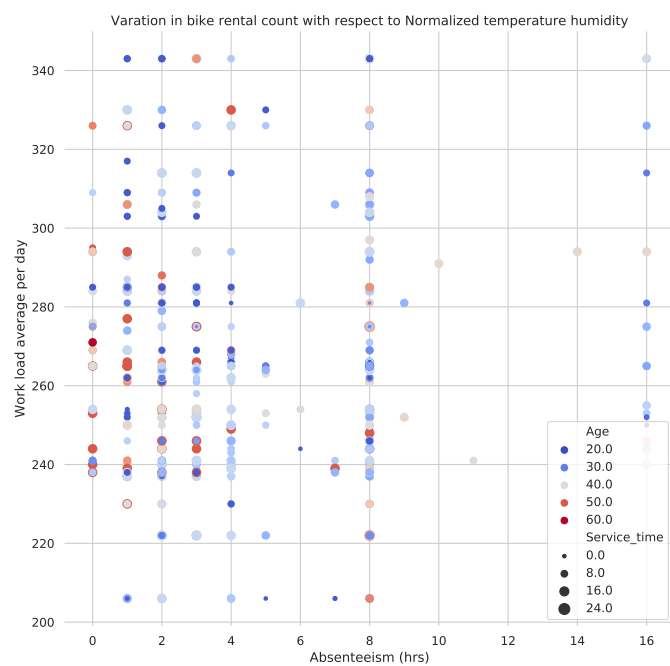


Figure 2.19: By plotting the work load per day vs absenteeism count with respect to service time and age. We can see that the employee having medium amount of work load i.e. 240 - 300 work load per day and who are below the age of 20 and having service time less than 8 years tend to be absent more often.



Figure 2.20: By plotting the absenteeism vs month of absenteeism with respect to education. It is clearly seen that employees having only a high school tend to be absent more often than the other employees.

---

# Feature Selection

---

To develop an efficient model, it is required to know which are the features which are of most importance in predicting the target variable. It is possible that many variable in our data set is less important in predicting the target than others or they may just be redundant. In this case it is required that we remove these variables from our data set to reduce its complexity. At times two features may carry the same information. In such case it is required to remove one of the features to avoid multi collinearity problems. There are many ways to perform feature selection or dimensionality reduction, but the method selected in this report is **Correlation Analysis** for continuous features and **ANOVA test** for categorical features.

## 3.0.1 Correlation Analysis

Correlation Analysis is a technique which helps to determine how strongly two features are related to each other (i.e. their co-variance). As the co-variance can vary from - infinity to + infinity, the correlation is used as it is a scaled version of the co-variance having values ranging from -1 to +1. A correlation plot is shown in figure 3.1. A correlation threshold of 0.8 is set and feature pairs of which exceeds this threshold, one of them is dropped. By examining table figure 3.1 it can be observed that features body mass index and weight are highly correlated. Hence the feature body mass index is dropped from the analysis.

## 3.0.2 Analysis Of Variance (ANOVA)

Analysis of Variance (ANOVA) is a statistical technique used to compute and compare the mean between two or more groups of observations. ANOVA makes use of two variables which are categorical variables and numeric variables of the data set. The python code in listing 3.1 is used to compute the p values of the feature and target variables. listing 3.1 also the output p- values to the threshold value of 0.05 and saves the feature names to be dropped in drop\_feat variable. The output of the code in listing 3.1 is shown in figure 3.2.

---

Listing 3.1: The python code used to compute the p values of the feature and target variables

---

```
1 ## ANOVA TEST FOR P VALUES
2 import statsmodels.api as sm
3 from statsmodels.formula.api import ols
```

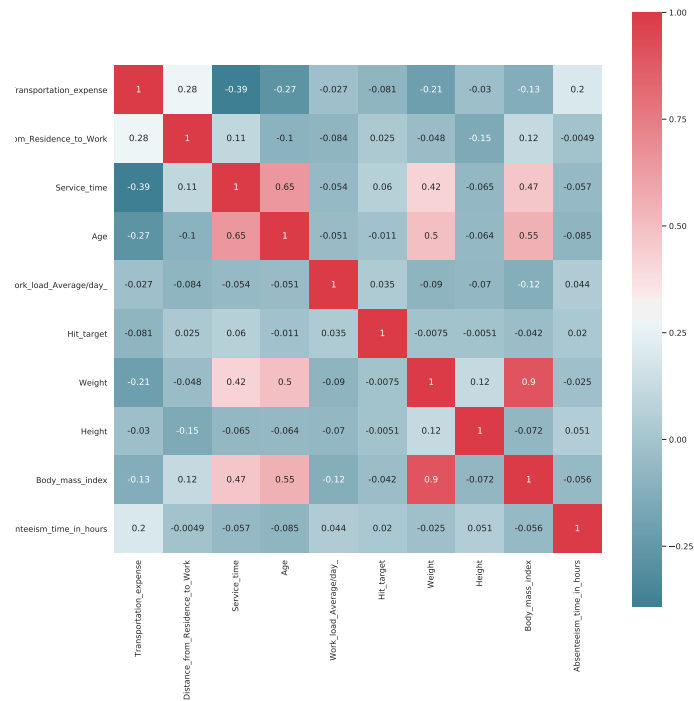


Figure 3.1: Correlation analysis of continuous variables.

```

4
5 anova_p = []
6 for i in obj_dtype:
7     buf = label + ' ~ ' + i
8     mod = ols(buf,data=df).fit()
9     anova_op = sm.stats.anova_lm(mod, typ=2)
10    print(anova_op)
11    anova_p.append(anova_op.iloc[0:1,3:4])
12    p = anova_op.loc[i,'PR(>F)']
13    if p >= 0.05:
14        drop_feat.append(i)

```

From figure 3.2 we can see that features Day\_of\_the\_week, Season, Social\_smoker and Education are greater than the p value threshold (0.05) and hence they are dropped along with Body\_mass\_index. Now the number of variables have been reduced from 21 to 16.

	sum_sq	df	F	PR(>F)
ID	1489.036365	35.0	4.146099	1.148030e-13
Residual	7223.875798	704.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Reason_for_absence	1397.477786	4.0	35.102023	7.515585e-27
Residual	7315.434376	735.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Month_of_absence	282.115342	11.0	2.214608	0.01228
Residual	8430.796821	728.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Day_of_the_week	67.201829	4.0	1.428262	0.222809
Residual	8645.710333	735.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Seasons	69.029361	3.0	1.959213	0.118713
Residual	8643.882801	736.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Disciplinary_failure	581.734796	1.0	52.799276	9.418497e-13
Residual	8131.177366	738.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Education	35.559279	3.0	1.005362	0.389784
Residual	8677.352883	736.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Son	339.960097	4.0	7.46065	0.000007
Residual	8372.952065	735.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Social_drinker	73.090287	1.0	6.243257	0.012683
Residual	8639.821875	738.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Social_smoker	23.110899	1.0	1.962743	0.161641
Residual	8689.801263	738.0	NaN	NaN
	sum_sq	df	F	PR(>F)
Pet	150.756818	5.0	2.584758	0.024951
Residual	8562.155344	734.0	NaN	NaN

Figure 3.2: Analysis Of Variance (ANOVA).

---

# Feature Engineering

---

It was noticed that the categories 1-21 in the feature `Reasons_for_absence` fall under one common theme which is Code of Diseases. Hence categories 1-21 are combined as a single category in the feature `Reason_for_absence`. In the same feature categories 22,23 and 28 relate to some kind of medical consultation and categories 24 and 25 relate to some kind of laboratory visit. Hence these categories are combined to form a single category for each respectively. Now the number of categories present in `Reason_for_absence` has reduced from 28 to 4.

### 4.0.1 Feature Scaling:

After the outlier analysis and other pre processing another skewness test is conducted on the data to check if they are skewed. As show in figure 4.1 `Absenteeism_time_in_hours` has a skew of about 1.1 which is much greater than rest of the features. In order to reduce the skewness of `Absenteeism_time_in_hours` a log transform is performed on `Absenteeism_time_in_hours`. This reduces the skewness of `Absenteeism_time_in_hours` to -0.23 as shown in figure 4.2.

After this the rest of the data is standardized in order to avoid any bias in the analysis results.

### 4.0.2 Creation of dummy variables

Now that most of the pre processing is over the last thing left to do is to convert the categorical variables into their respective dummy variables but creating a feature for each category group in the categorical feature. As our categorical variables are in the form of numbers and not strings we need not convert them to numbers, but it is required that we create dummy variables for each categorical group. As the categorical features in our data set are nominal categorical data and not ordinal categories the model shouldn't give higher precedence to a higher number. For example, the categories present in the feature `seasons` is 1, 2, 3 and 4 which corresponds to spring, summer, fall and winter respectively. As the categories are numeric, the model would give a higher precedence to 4 as it is numerically greater than the rest. Which is not so as they are not ordinal categories and each number refers to a specific season. Hence each group is converted to a binary category. Where 1 imply the presence of that group and 0 represents absence. To make it less complex the first dummy variable is dropped, as when all the other groups are absent it imply that the dropped group is present. This is achieved by running the command in listing 4.1.



```

*****
statistical properties of Transportation_expense:
DescribeResult(nobs=740, minmax=(118.0, 378.0), mean=220.62567567567567, variance=4373.354955564495, skewness=0.3756674357536574, kurtosis=-0.33770597431634375)
*****
statistical properties of Distance_from_Residence_to_Work:
DescribeResult(nobs=740, minmax=(5.0, 52.0), mean=29.63108108108108, variance=220.13029111655638, skewness=0.3114498278842765, kurtosis=-1.261270969913851)
*****
statistical properties of Service_time:
DescribeResult(nobs=740, minmax=(1.0, 24.0), mean=12.445945945945946, variance=17.386790037669606, skewness=-0.34301872345212286, kurtosis=-0.19060374758858512)
*****
statistical properties of Age :
DescribeResult(nobs=740, minmax=(27.0, 53.0), mean=36.14054054054054, variance=37.37129064111472, skewness=0.4894588639031262, kurtosis=-0.25563936337912363)
*****
statistical properties of Work_load Average/day_:
DescribeResult(nobs=740, minmax=(206.0, 343.0), mean=266.94594594594594, variance=1026.0106060051933, skewness=0.5610733905805646, kurtosis=-0.18089326615469226)
*****
statistical properties of Hit_target:
DescribeResult(nobs=740, minmax=(87.0, 100.0), mean=94.92297297297297, variance=9.47849723878141, skewness=-0.44564951370083833, kurtosis=-0.3920431683313459)
*****
statistical properties of Weight:
DescribeResult(nobs=740, minmax=(56.0, 108.0), mean=79.03513513513514, variance=165.9771129722415, skewness=0.016966890485138695, kurtosis=-0.9158612723351429)
*****
statistical properties of Height:
DescribeResult(nobs=740, minmax=(165.0, 175.0), mean=170.20675675675676, variance=3.7014427824306027, skewness=-0.44993040994803707, kurtosis=0.7658511236420802)
*****
statistical properties of Body_mass index:
DescribeResult(nobs=740, minmax=(19.0, 38.0), mean=26.681081081081082, variance=18.374472442672715, skewness=0.3001649161585004, kurtosis=-0.3226064788195093)
*****
statistical properties of Absenteeism_time_in_hours:
DescribeResult(nobs=740, minmax=(0.0, 16.0), mean=4.425675675675675, variance=11.790138243791832, skewness=1.103093333423475, kurtosis=1.2887445861552127)
*****

```

Figure 4.1: Skewtest after outlier analysis shows us that the target Absenteeism\_time\_in\_hours still has a large skew of 1.1 when compared to the other features.

Listing 4.1: Creation of dummy variables.

```

1
2  ## Creating dummy variables
3  # Where x is all the features in our analysis.
4
5  X_lr = pd.get_dummies(X, columns = obj_dtype, drop_first = True)

```

```

*****
statistical properties of Transportation_expense:
DescribeResult(nobs=740, minmax=(118.0, 378.0), mean=220.62567567567567, variance=4373.354955564495, skewness=0.3756674357536574, kurtosis=-0.33770597431634375)
*****
statistical properties of Distance_from Residence to Work:
DescribeResult(nobs=740, minmax=(5.0, 52.0), mean=29.63108108108108, variance=220.13029111655638, skewness=0.3114498278842765, kurtosis=-1.261270969913851)
*****
statistical properties of Service_time:
DescribeResult(nobs=740, minmax=(1.0, 24.0), mean=12.445945945945946, variance=17.386790037669606, skewness=-0.34301872345212286, kurtosis=-0.19060374758858512)
*****
statistical properties of Age :
DescribeResult(nobs=740, minmax=(27.0, 53.0), mean=36.14054054054054, variance=37.37129064111472, skewness=0.4894588639031262, kurtosis=-0.25563936337912363)
*****
statistical properties of Work_load Average/day_:
DescribeResult(nobs=740, minmax=(206.0, 343.0), mean=266.94594594594594, variance=1026.0106060051933, skewness=0.5610733905805646, kurtosis=-0.18089326615469226)
*****
statistical properties of Hit target:
DescribeResult(nobs=740, minmax=(87.0, 100.0), mean=94.92297297297297, variance=9.47849723878141, skewness=-0.44564951370083833, kurtosis=-0.3920431683313459)
*****
statistical properties of Weight:
DescribeResult(nobs=740, minmax=(56.0, 108.0), mean=79.03513513513514, variance=165.9771129722415, skewness=0.016966890485138695, kurtosis=-0.9158612723351429)
*****
statistical properties of Height:
DescribeResult(nobs=740, minmax=(165.0, 175.0), mean=170.20675675675676, variance=3.7014427824306027, skewness=-0.44993040994803707, kurtosis=0.7658511236420802)
*****
statistical properties of Body_mass_index:
DescribeResult(nobs=740, minmax=(19.0, 38.0), mean=26.681081081081082, variance=18.374472442672715, skewness=0.3001649161585004, kurtosis=-0.3226064788195093)
*****
statistical properties of Absenteeism_time_in_hours:
DescribeResult(nobs=740, minmax=(0.0, 2.833213344056216), mean=1.4873456789490085, variance=0.4367702701986841, skewness=-0.23144444623756816, kurtosis=-0.4684714403096022)
*****

```

Figure 4.2: Skew test after performing log transform on Absenteeism\_time\_in\_hours that the skewness dropped to -0.23 for Absenteeism\_time\_in\_hours.

---

# Modeling

---

## 5.1 Model Selection

It has been noted in previous stages of our analysis that for different combinations of the independent variables, the count is different. It can be seen that the dependent variable is a continuous variable and hence the type of model that would be developed for this problem is a regression model. It can also be seen that there are many features which may contribute to our regression model. That would make this a multivariate regression problem. However the data is not normally distributed, hence if we are not able to generate a model with good performance we may bin the target variable and perform multi class classification on it.

## 5.2 Methodology:

Model Evaluation is an integral part of the model development process as it helps us find the best model for representing our data. It also helps to evaluate as to how it would perform on new data. In order to develop an efficient and accurate model to predict our target variable we shall use a combination of three different methods. The three different methods used in our analysis is given below.

- Hold-Out Method
- Hyper parameter Tuning
- Cross-Validation Technique

### 5.2.1 Hold-Out Method

As evaluating model performance on training data set may lead to develop an over fitted model. Due to this it is required to test the model on a separate data set. Hence the original data set is split into training and testing data. The training data set is used to build a predictive model and the testing data is used to evaluate the model performance.

### 5.2.2 Hyper parameter Tunning

Hyper parameter Tuning is a method in which the parameters of the model is to be set before training. Scikit-Learn implements a set of sensible default hyper parameters but it may not be optimal. In our analysis we shall use two types of hyper parameter tuning methodology which are *Random Search CV* and *Grid Search CV*. The major difference between them is the run time. As randomized search is drastically lower than grid search. Random search is faster than grid search as we do not provide a set of discrete values to be searched. Rather we provide a statistical distribution for each parameter from which values may be randomly sampled. Where as in grid search a discrete set of values are provided for each parameter and several models are built on each of its combinations which makes it very laboursome.

### 5.2.3 K-Fold Cross Validation Technique

In K fold cross validation technique the data is divided into k subsets of equal size. We build the model K times, each time leaving out a subset from training. This sub set which was left out will be used for testing. This method helps us develop an unbiased estimate of the model performance.

## 5.3 Regression Model Building

An ordinary least square linear regression model is first built on the entire data, in order to check for multicollinearity. Before the model is built in multivariate linear regression it is required to do some processing on the data before hand. The equation around which the multivariate regression in statsmodels works is shown below.

$$y = \beta_0 + \beta_1 * X_1 + \beta_2 * X_2 + .... + \beta_n * X_n \quad (5.1)$$

Where  $\beta_1, \beta_2$  till  $\beta_n$  are the coefficients for columns  $X_1, X_2$  till  $X_n$  respectively and  $\beta_0$  represents the intercept. Unfortunately in statsmodel the  $\beta_0$  column is not added by default to our model and hence we are required to add a column with a constant value of one to the data set. Now that this column is added and the dummy variables are created we can go ahead with building the model. It can be noted from the summary of the model that almost none of the features are significant. This could be due to multicollinearity present in the model. Hence backward elimination is performed in order to remove the multicollinearity present in the model.

### 5.3.1 Backward Elimination:

In backward elimination, we start with all the features and remove the least significant feature at each iteration which improves the performance of the model. In this procedure we keep an eye on the adjusted r squared value. As we remove insignificant features the model performance increases and the adjusted r square also increases. The iteration is stopped when the value of adjusted r squared drops when compared to the previous iteration. When the adjusted r squared dropped it means that all the features are statistically significant for our model and removal of a single feature beyond this point would reduce the performance of our model. We start the backward elimination process with a total of 69 features with a r square of 0.52 and adjusted r squared value of 0.488.

After 31 iterations it can be noted that the adj r squared has risen to 0.99 and on the 32<sup>nd</sup> the adjusted R squared value drops to 0.498. Therefore we stop the backward elimination at the 31<sup>st</sup> iteration. After conducting backward elimination it can be seen that the number of features have dropped from 69 to 37.

### Multivariate Linear Regression Model after backward elimination

We Start building our model by using the simplest model to the most complex model. Therefore we start our model building using the multiple linear regression model. It can be seen from figure 5.1 that the **Adjusted R-squared is 0.41** which is not impressive, as **it means 41% of the variance of the data is explained by the model**. The code to develop the model and its output can be seen in figure 5.1.

#### Multivariate linear regression

```
In [1351]: from sklearn.cross_validation import train_test_split
           #Divide data into train and test
           train, test = train_test_split(df, test_size=0.2)

In [1352]: # Train the model using the training sets
           model = sm.OLS(train.iloc[:,36], train.iloc[:,0:36]).fit()

           # make the predictions by the model
           predictions_LR = model.predict(test.iloc[:,0:36])

           r2 = r2_score(test.iloc[:,36], predictions_LR)
           mse = mean_squared_error(test.iloc[:,36], predictions_LR)

           print('Linear Regression Model Performance:')
           print('R-squared = {:.2}'.format(r2))
           print('MSE = ',mse)

Linear Regression Model Performance:
R-squared = 0.41.
MSE = 0.24742723773384243
```

Figure 5.1: Multivariate linear regression model after backward elimination with  $R^2$  value 0.41 and MSE 0.24.

### Decision Tree Regressor after backward elimination

Now a decision tree regressor model is built on to our data set. At first a model is built on the default parameters of the Decision tree regressor. The default values for the parameters controlling the size of the trees such as `max_depth`, `min_samples_leaf`, etc lead to fully grown and unpruned trees which could require large memory consumption and increase the model complexity. It can also cause overfitting of the model. To reduce the complexity and memory consumption the size of the trees should be controlled by setting those parameter values. In our analysis we are going to tune the `max_depth` parameter which controls the maximum depth of the tree. The default model without any tuning is taken as the base line model.

After this a grid of intervals of 5 is taken from 5 - 50 for `max_depth` are given to random search cv. The k fold cross validation parameter given is K is 5 and `n_iter` is 5, which the cross validation is to be repeated 5 times. The optimal depth from random search cv is 10. Now that we narrowed down our search criteria, grid search cv can be conducted on range 5 - 20 `max_depth` with an interval of 2.

Here it was found that the best parameters for max\_depth is 7. The code to develop the models and their output can be seen in figure 5.2, 5.3 & 5.4.

### Decision tree

```
In [464]: #dropping column b0
train = train.drop(["b0"],axis = 1)
test = test.drop(["b0"],axis = 1)

In [465]: #Decision Regressor

from sklearn.tree import DecisionTreeRegressor

X_train = train.iloc[:,0:35]
y_train = train.iloc[:,35]

X_test = test.iloc[:,0:35]
y_test = test.iloc[:,35]

DT_reg = DecisionTreeRegressor(random_state=0).fit(X_train,y_train)

DT_reg.get_params()

#Apply model on test data
predictions_DT = DT_reg.predict(X_test)

#R^2
DT_r2 = r2_score(y_test, predictions_DT)
#Calculating MSE
DT_mse = mean_squared_error(y_test, predictions_DT)

print('Decision Regressor Model Performance:')
print('Default Parameters = ',DT_reg.get_params())
print('R-squared = {:.0.2}'.format(DT_r2))
print('MSE = ',DT_mse)

Decision Regressor Model Performance:
Default Parameters = {'criterion': 'mse', 'max_depth': None, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_weight_fraction_leaf': 0.0, 'presort': False, 'random_state': 0, 'splitter': 'best'}
R-squared = 0.038.
MSE = 0.4015202539881518
```

Figure 5.2: Default Decision tree model with  $R^2$  value 0.038 and MSE 0.4

### Random Forest Regressor after backward elimination

Random forest is an ensemble that consists of many decision trees. At first a model is built on the default parameters of the Random Forest regressor. To reduce the complexity and memory consumption the max\_depth and n\_estimators parameters are tuned. n\_estimators tell the amount of trees to be used in the model. The default model without any tuning is taken as the base line model.

After this a grid of intervals of 2 is taken from 1 - 20 max\_depth and a grid of intervals of 2 is taken from 1 - 100 for n\_estimators given to random search cv. The k fold cross validation parameter given is K is 5 and n\_iter is 5, which the cross validation is to be repeated 5 times. The optimal depth from random search cv is 9 and n\_estimators is 15. Now that we narrowed down our search criteria, grid search cv can be conducted on 5 - 15 for max\_depth and intervals of 1 is taken from 11 - 20 for n\_estimators. Here it was found that the best parameters for max\_depth is 12 and n\_estimators is 16. The code to develop the models and their output can be seen in figure 5.5, 5.6 & 5.7.

```

In [466]: ##Random Search CV
from sklearn.model_selection import RandomizedSearchCV
np.random.seed(0)
RDT = DecisionTreeRegressor(random_state = 0)
depth = list(range(5,50,5))
# Create the random grid
randDT_grid = {'max_depth': depth}

randomcv_DT = RandomizedSearchCV(RDT, param_distributions = randDT_grid, n_iter = 5, cv = 5, random_state = 0)
randomcv_DT = randomcv_DT.fit(X_train,y_train)

predictions_RDT = randomcv_DT.predict(X_test)

view_best_params_RDT = randomcv_DT.best_params_

best_model = randomcv_DT.best_estimator_

predictions_RDT = best_model.predict(X_test)

#R^2
RDT_r2 = r2_score(y_test, predictions_RDT)
#Calculate MSE
RDT_mse = mean_squared_error(y_test, predictions_RDT)

print('Random Search CV Decision Regressor Model Performance:')
print('Best Parameters = ',view_best_params_RDT)
print('R-squared = {:.0.2}'.format(RDT_r2))
print('MSE = ',RDT_mse)
print('*****')

Random Search CV Decision Regressor Model Performance:
Best Parameters = {'max_depth': 10}
R-squared = 0.21.
MSE = 0.32953047191433227
*****

```

Figure 5.3: Random search CV Decision tree model with a depth of 10, having  $R^2$  value 0.21 and MSE 0.32.

### Gradient Boosting after backward elimination

Gradient boosting is an ensemble boosting method in which a collection of regressors are built in series. It is a method of converting a sequence of weak learners to a complex model. Each regressor's prediction is based on the previous regressors by adding weights accordingly. We shall now implement the Gradient boosting model on our data. At first a model is built on the default parameters of the Random Forest regressor. To reduce the complexity and memory consumption the `max_depth` and `n_estimators` parameters are tuned. `n_estimators` tells the amount of trees to be used in the model. The default model without any tuning is taken as the base line model.

After this a grid of intervals of 2 is taken from 1 - 10 `max_depth` and a grid of intervals of 10 is taken from 50 - 150 for `n_estimators` given to random search cv. The k fold cross validation parameter given is K is 5 and `n_iter` is 5, which the cross validation is to be repeated 5 times. The optimal depth from random search cv is 3 and `n_estimators` is 50. Now that we narrowed down our search criteria, grid search cv can be conducted on range 1 - 5 for `max_depth` and intervals of 5 is taken from 40 - 80 for `n_estimators`. Here it was found that the best parameters for `max_depth` is 4 and `n_estimators` is 45. The code to develop the models and their output can be seen in figure 5.8, 5.9 & 5.10.



```

In [470]: ##Grid Search CV

from sklearn.model_selection import GridSearchCV

Gridregr = DecisionTreeRegressor(random_state = 0)
depth = list(range(1,10,1))

# Create the grid
grid_search = {'max_depth': depth}

## Grid Search Cross-Validation with 5 fold CV
gridcv_GDT = GridSearchCV(Gridregr, param_grid = grid_search, cv = 5)
gridcv_GDT = gridcv_GDT.fit(X_train,y_train)
view_best_params_GDT = gridcv_GDT.best_params_

#Apply model on test data
predictions_GDT = gridcv_GDT.predict(X_test)

#R^2
GDT_r2 = r2_score(y_test, predictions_GDT)
#Calculate MSE
GDT_mse = mean_squared_error(y_test, predictions_GDT)
#Calculate MAPE
GDT_mape = MAPE(y_test, predictions_GDT)

print('Grid Search CV Decision Regressor Model Performance:')
print('Best Parameters = ',view_best_params_GDT)
print('R-squared = {:.0.2}'.format(GDT_r2))
print('MSE = ',(GDT_mse))
print('*****')

Grid Search CV Decision Regressor Model Performance:
Best Parameters = {'max_depth': 7}
R-squared = 0.18.
MSE = 0.34127194268942496
*****

```

Figure 5.4: Grid search CV Decision tree model with a depth of 7, having  $R^2$  value 0.18 and MSE 0.34.

```

In [477]: from sklearn.ensemble import RandomForestRegressor

RF_reg = RandomForestRegressor(random_state=0).fit(X_train,y_train)
RF_reg.get_params()

#Apply model on test data
predictions_RF = RF_reg.predict(X_test)

#R^2
RF_r2 = r2_score(y_test, predictions_RF)
#Calculating MSE
RF_mse = np.mean(( y_test - predictions_RF)**2)

print('Random Forest Regressor Model Performance:')
print('Default Parameters = ',RF_reg.get_params())
print('R-squared = {:.0.2}'.format(RF_r2))
print('MSE = ',RF_mse)
print('*****')

Random Forest Regressor Model Performance:
Default Parameters = {'bootstrap': True, 'criterion': 'mse', 'max_depth': None, 'max_features': 'auto',
x_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'mi
amples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 10, 'n_jobs': 1, 'oob_score': False, '
om_state': 0, 'verbose': 0, 'warm_start': False}
R-squared = 0.23.
MSE = 0.3230720690534794
*****

```

Figure 5.5: Default Random forest model with  $R^2$  value 0.23 and MSE 0.32



```

In [473]: ##Random Search CV
from sklearn.model_selection import RandomizedSearchCV

RRF = RandomForestRegressor(random_state = 0)
n_estimator = list(range(1,20,2))
depth = list(range(1,100,2))

# Create the random grid
rand_grid = {'n_estimators': n_estimator,
             'max_depth': depth}

randomcv_rf = RandomizedSearchCV(RRF, param_distributions = rand_grid, n_iter = 5, cv = 5, random_state = 0)
randomcv_rf = randomcv_rf.fit(X_train,y_train)
predictions_RRF = randomcv_rf.predict(X_test)

view_best_params_RRF = randomcv_rf.best_params_

best_model = randomcv_rf.best_estimator_

predictions_RRF = best_model.predict(X_test)

#R^2
RRF_r2 = r2_score(y_test, predictions_RRF)
#Calculating MSE
RRF_mse = np.mean(( y_test - predictions_RRF)**2)

print('Random Search CV Random Forest Regressor Model Performance:')
print('Best Parameters = ',view_best_params_RRF)
print('R-squared = {:.2}'.format(RRF_r2))
print('MSE = ',RRF_mse)

Random Search CV Random Forest Regressor Model Performance:
Best Parameters = {'n_estimators': 15, 'max_depth': 9}
R-squared = 0.29.
MSE = 0.29844963753444115

```

Figure 5.6: Random search CV Random forest model with a depth of 9 and n estimators 15. Having  $R^2$  value 0.29 and MSE 0.29.

### 5.3.2 Principal component analysis:

Since non of the above models give us a suitable rsquared value. We now perform PCA on our data to reduce the complexity of the data and to make sure that all the variables are independent of each other. Principal component analysis is a technique for feature extraction. In which our data is combined in a specific way, then we can drop the “least important” variables while still retaining the most valuable parts of all of the variables. After performing PCA on the data we can see that the First Principal component explains about 35% of the variance and 20 principal components explain 98.01% of the variance. This can be seen in figure 5.11 & 5.12

#### Multivariate Linear Regression Model after backward elimination and PCA

We Start building our model by using 20 principal components which explains 98% of the variance of the target variable. It can be seen from figure 5.13 that the **Adjusted R-squared is 0.28** which is not impressive, as **it means 28% of the variance of the data is explained by the model**. The code to develop the model and its output can be seen in figure 5.13.

#### Decision Tree Regressor after backward elemination and PCA

Now a decision tree regressor model is built on to our data set. At first a model is built on the default parameters of the Decision tree regressor. After this a grid of intervals of 5 is taken from 5 - 50 for

```

In [476]: ## Grid Search CV
from sklearn.model_selection import GridSearchCV

regr = RandomForestRegressor(random_state = 0)
n_estimator = list(range(11,20,1))
depth = list(range(5,15,2))

# Create the grid
grid_search = {'n_estimators': n_estimator,
               'max_depth': depth}

## Grid Search Cross-Validation with 5 fold CV
gridcv_rf = GridSearchCV(regr, param_grid = grid_search, cv = 5)
gridcv_rf = gridcv_rf.fit(X_train,y_train)
view_best_params_GRF = gridcv_rf.best_params_

#Apply model on test data
predictions_GRF = gridcv_rf.predict(X_test)

#R^2
GRF_r2 = r2_score(y_test, predictions_GRF)
#Calculating MSE
GRF_mse = np.mean(( y_test - predictions_GRF)**2)

print('Grid Search CV Random Forest Regressor Model Performance:')
print('Best Parameters = ',view_best_params_GRF)
print('R-squared = {:.2}'.format(GRF_r2))
print('MSE = ',(GRF_mse))

Grid Search CV Random Forest Regressor Model Performance:
Best Parameters = {'max_depth': 7, 'n_estimators': 18}
R-squared = 0.29.
MSE = 0.2979419859804356

```

Figure 5.7: Grid search CV Random forest model with a depth of 7 n estimators 18. Having  $R^2$  value 0.29 and MSE 0.29.

max\_depth are given to random search cv. The k fold cross validation parameter given is K is 5 and n\_iter is 5, which the cross validation is to be repeated 5 times. The optimal depth from random search cv is 10. Now that we narrowed down our search criteria, grid search cv can be conducted on range 1 - 15 max\_depth with an interval of 1. Here it was found that the best parameters for max\_depth is 5. The code to develop the models and their output can be seen in figure 5.14, 5.15 & 5.16.

### Random Forest Regressor after backward elimination and PCA

Random forest is an ensemble that consists of many decision trees At first a model is built on the default parameters of the Random Forest regressor. To reduce the complexity and memory consumption the max\_depth and n\_estimators parameters is tuned n\_estimators tell the amount of trees to be used in the model. The default model without any tuning is taken as the base line model.

After this a grid of intervals of 2 is taken from 1 - 20 max\_depth and a grid of intervals of 2 is taken from 1 - 100 for n\_estimators given to random search cv. The k fold cross validation parameter given is K is 5 and n\_iter is 5, which the cross validation is to be repeated 5 times. The optimal depth from random search cv is 9 and n\_estimators is 15. Now that we narrowed down our search criteria, grid search cv can be conducted on 5 - 15 for max\_depth and intervals of 1 is taken from 11 - 20 for n\_estimators. Here it was found that the best parameters for max\_depth is 12 and n\_estimators is 16. The code to develop the models and their output can be seen in figure 5.17, 5.18

## Gradient Boost

```
In [479]: #Gradient Boost
from sklearn.ensemble import GradientBoostingRegressor

gbt = GradientBoostingRegressor(random_state= 0).fit(X_train,y_train)

predictions_gbt = gbt.predict(X_test)

gbt.get_params()

#R^2
GBR_r2 = r2_score(y_test, predictions_gbt)
#Calculate MSE
GBR_mse = mean_squared_error(y_test, predictions_gbt)

print('Gradient Boosting Regressor Model Performance:')
print('Default Parameters = ',gbt.get_params())
print('R-squared = {:.2}'.format(GBR_r2))
print('MSE = ',GBR_mse)
print('*****')

Gradient Boosting Regressor Model Performance:
Default Parameters = {'alpha': 0.9, 'criterion': 'friedman_mse', 'init': None, 'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 3, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'presort': 'auto', 'random_state': 0, 'subsample': 1.0, 'verbose': 0, 'warm_start': False}
R-squared = 0.28.
MSE = 0.3004434155173034
*****
```

Figure 5.8: Default Gradient Boosting model with max depth 3 and n estimator 100, has a  $R^2$  value 0.28 and MSE 0.3

& 5.19.

## Gradient Boosting after backward elimination and PCA

Gradient boosting is an ensemble boosting method in which a collection of regressors are built in series. It is a method of converting a sequence of weak learners to a complex model. Each regressor's prediction is based on the previous regressors by adding weights accordingly. We shall now implement the Gradient boosting model on our data. At first a model is built on the default parameters of the Random Forest regressor. To reduce the complexity and memory consumption the max\_depth and n\_estimators parameters are tuned. n\_estimators tell the amount of trees to be used in the model. The default model without any tuning is taken as the base line model.

After this a grid of intervals of 2 is taken from 1 - 10 max\_depth and a grid of intervals of 10 is taken from 50 - 150 for n\_estimators given to random search cv. The k fold cross validation parameter given is K is 5 and n\_iter is 5, which the cross validation is to be repeated 5 times. The optimal depth from random search cv is 3 and n\_estimators is 50. Now that we narrowed down our search criteria, grid search cv can be conducted on range 1 - 5 for max\_depth and intervals of 5 is taken from 40 - 80 for n\_estimators. Here it was found that the best parameters for max\_depth is 4 and n\_estimators is 45. The code to develop the models and their output can be seen in figure 5.20, 5.21 & 5.22.

## Random Search CV Gradient boosting

```
In [483]: ##Random Search CV
rGBR = GradientBoostingRegressor(random_state = 0)
#loss = ['ls','lad','huber','quantile']
n_estimator = list(range(50,150,10))
#max_feat = ['auto','sqrt','log2']
depth = list(range(1,10,2))

# Create the random grid
rand_GBT = {'loss': loss,
            'n_estimators': n_estimator,
            'max_features': max_feat,
            'max_depth': depth}

randomcv_gbt = RandomizedSearchCV(rGBR, param_distributions = rand_GBT, n_iter = 5, cv = 5, random_state=0)
randomcv_gbt.fit(X_train,y_train)
predictions_GBT = randomcv_gbt.predict(X_test)

view_best_params_GBT = randomcv_gbt.best_params_

#R^2
rGBR_r2 = r2_score(y_test, predictions_GBT)
#Calculate MSE
rGBR_mse = mean_squared_error(y_test, predictions_GBT)

print('Random Search CV Gradient Boosting Regressor Model Performance:')
print('Best Parameters = ',view_best_params_GBT)
print('R-squared = {:.2}'.format(rGBR_r2))
print('MSE = ',rGBR_mse)
print('*****')

Random Search CV Gradient Boosting Regressor Model Performance:
Best Parameters = {'n_estimators': 50, 'max_depth': 3}
R-squared = 0.25.
MSE = 0.31125407863782906
*****
```

Figure 5.9: Random search CV Gradient Boosting model with a depth of 3 and n estimators 50. Having  $R^2$  value 0.25 and MSE 0.31.

## 5.4 Classification Model Building

The performance parameters of the models built in the previous section was below par. This could be because the variables are not nominally distributed. Hence the target variable is binned into 4 categories and multi class classification is performed on it. The bin ranges for the classification is such that values absenteeism time in hours of 0 considered as less than one hour as category 1, values absenteeism time in hours of 1 considered as one hour as category 2, the values absenteeism time in hours of 2-5 considered as half a day of absenteeism as category 3 and the values with absenteeism time in hours of 5-16 is considered as a whole day of absenteeism as category 4, the distribution of the new target variable is shown in figure 5.23

Before building models we conduct chi squared test on the categorical data for feature selection this is shown in figure 5.24. and the normal procedure of scaling is conducted on the data set. According to the chi squared test Day\_of\_the\_week education and pet is dropped along with weight from the correlation analysis.

```

In [484]: ## Grid Search CV

gGBR = GradientBoostingRegressor(random_state=0)
#loss = ['ls', 'lad', 'huber', 'quantile']
n_estimator = list(range(40,80,5))
#max_feat = ['auto', 'sqrt', 'log2']
depth = list(range(1,5,1))

# Create the random grid
grid_GBT = {'loss': loss,
            'n_estimators': n_estimator,
            #'max_features': max_feat,
            'max_depth': depth}

## Grid Search Cross-Validation with 5 fold CV
gridcv_GBT = GridSearchCV(gGBR, param_grid = grid_GBT, cv = 5)
gridcv_GBT = gridcv_GBT.fit(X_train,y_train)
view_best_params_gridGRF = gridcv_GBT.best_params_

#Apply model on test data
predictions_gridGBT = gridcv_GBT.predict(X_test)

#R^2
gGBR_r2 = r2_score(y_test, predictions_gridGBT)
#Calculate MSE
gGBR_mse = mean_squared_error(y_test, predictions_gridGBT)

#Calculate MAPE
gGBR_mape = MAPE(y_test, predictions_gridGBT)

print('Grid Search CV Gradient Boosting Regressor Model Performance:')
print('Best Parameters = ',view_best_params_gridGRF)
print('R-squared = {:.2}'.format(gGBR_r2))
print('MSE = ', gGBR_mse)
print('*****')

Grid Search CV Gradient Boosting Regressor Model Performance:
Best Parameters = {'max_depth': 4, 'n_estimators': 45}
R-squared = 0.27.
MSE = 0.3065385845417901
*****

```

Figure 5.10: Grid search CV Gradient Boosting model with a depth of 4 n estimators 45. Having  $R^2$  value 0.27 and MSE 0.3.

### Decision tree multiclass classifier

A multiclass decision tree classifier is built on the train data to predict as to predict the class of absenteeism. The code confusion matrix is shown in figure 5.25. The decision tree classifier has an accuracy of 59.7%.

### Logistic regression multi class classifier

A Logistic regression multi class classifier model is built on the train data to predict as to predict the class of absenteeism. The code and output is shown in figure 5.26. The Logistic regression classifier has an accuracy of 61.6%.

### Random forest multi class classifier

A Random forest multi class classifier model is built on the train data to predict as to predict the class of absenteeism. The code and output is shown in figure 5.27. The classifier has an accuracy of 61.6%.

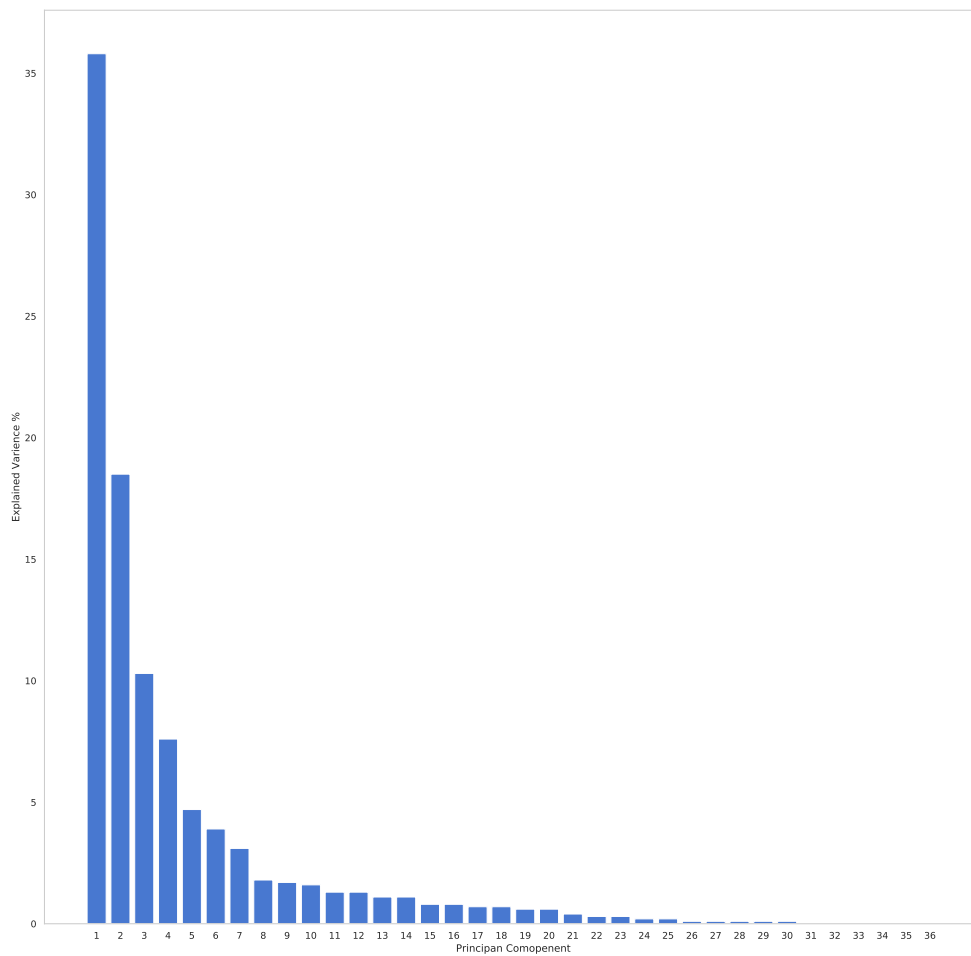


Figure 5.11: Explained variance of each principal component.

### KNN classifier

A KNN classifier model is built on the train data to predict as to predict the class of absenteeism. The code and output is shown in figure 5.28. The KNN classifier has an accuracy of 50.6%.

### Naive Bayes

Naive Bayes classifier model is built on the train data to predict as to predict the class of absenteeism. The code and output is shown in figure 5.29. The Naive Bayes classifier has an accuracy of 47.3%.s

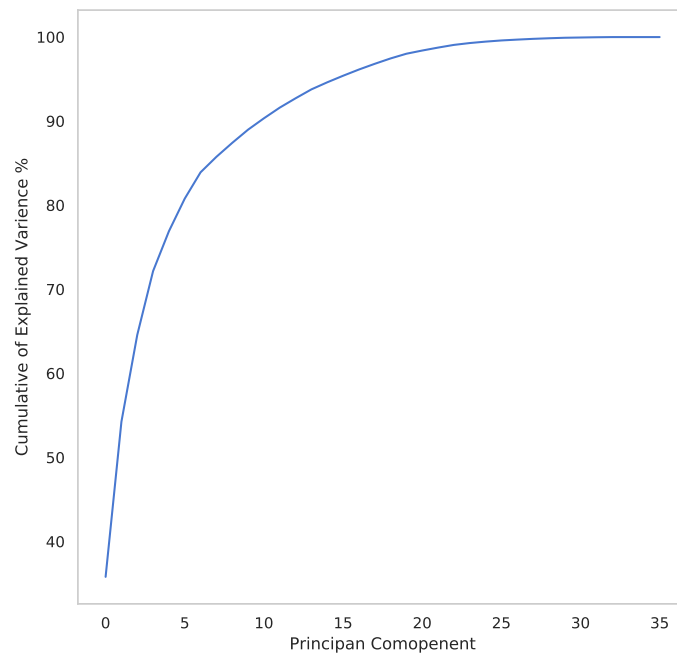


Figure 5.12: Cumulative of explained variance of PCA.

### Multivariant linear regression

```
In [1351]: from sklearn.cross_validation import train_test_split
           #Divide data into train and test
           train, test = train_test_split(df, test_size=0.2)

In [1352]: # Train the model using the training sets
           model = sm.OLS(train.iloc[:,36], train.iloc[:,0:36]).fit()

           # make the predictions by the model
           predictions_LR = model.predict(test.iloc[:,0:36])

           r2 = r2_score(test.iloc[:,36], predictions_LR)
           mse = mean_squared_error(test.iloc[:,36], predictions_LR)

           print('Linear Regression Model Performance:')
           print('R-squared = {:.2}'.format(r2))
           print('MSE = ',mse)

Linear Regression Model Performance:
R-squared = 0.41.
MSE = 0.24742723773384243
```

Figure 5.13: Multivariant linear regression model after backword elimination and PCA with  $R^2$  value 0.28 and MSE 0.29.

```

In [507]: #Decision Regressor

from sklearn.tree import DecisionTreeRegressor

DT_reg = DecisionTreeRegressor(max_depth = 2, random_state=0).fit(X_train,y_train)

DT_reg.get_params()

#Apply model on test data
predictions_DT = DT_reg.predict(X_test)

#R^2
DT_r2 = r2_score(y_test, predictions_DT)
#Calculating MSE
DT_mse = mean_squared_error(y_test, predictions_DT)

print('Decision Regressor Model Performance:')
print('Default Parameters = ',DT_reg.get_params())
print('R-squared = {:.4}'.format(DT_r2))
print('MSE = ',DT_mse)
print('*****')

Decision Regressor Model Performance:
Default Parameters = {'criterion': 'mse', 'max_depth': 2, 'max_features': None, 'max_leaf_nodes': None,
n_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'm
eight_fraction_leaf': 0.0, 'presort': False, 'random_state': 0, 'splitter': 'best'}
R-squared = 0.1032.
MSE = 0.3744484881210333
*****

```

Figure 5.14: Default Decision tree model with  $R^2$  value 0.10 and MSE 0.37



### Decision tree Random Search CV

```
In [508]: ##Random Search CV
from sklearn.model_selection import RandomizedSearchCV
np.random.seed(0)
RDT = DecisionTreeRegressor(random_state = 0)
depth = list(range(5,50,5))
# Create the random grid
randDT_grid = {'max_depth': depth}

randomcv_DT = RandomizedSearchCV(RDT, param_distributions = randDT_grid, n_iter = 5, cv = 5, random_state = 0)
randomcv_DT = randomcv_DT.fit(X_train,y_train)

predictions_RDT = randomcv_DT.predict(X_test)

view_best_params_RDT = randomcv_DT.best_params_

best_model = randomcv_DT.best_estimator_

predictions_RDT = best_model.predict(X_test)

#R^2
RDT_r2 = r2_score(y_test, predictions_RDT)
#Calculate MSE
RDT_mse = mean_squared_error(y_test, predictions_RDT)

print('Random Search CV Decision Regressor Model Performance:')
print('Best Parameters = ',view_best_params_RDT)
print('R-squared = {:.2}'.format(RDT_r2))
print('MSE = ',RDT_mse)
print('*****')

Random Search CV Decision Regressor Model Performance:
Best Parameters = {'max_depth': 10}
R-squared = 0.015.
MSE = 0.41137102943099124
*****
```

Figure 5.15: Random search CV Decision tree model with a depth of 10, having  $R^2$  value 0.015 and MSE 0.41.

### Decision tree Grid Search CV

```
In [513]: ##Grid Search CV

from sklearn.model_selection import GridSearchCV

Gridregr = DecisionTreeRegressor(random_state = 0)
depth = list(range(1,15,1))

# Create the grid
grid_search = {'max_depth': depth}

## Grid Search Cross-Validation with 5 fold CV
gridcv_GDT = GridSearchCV(Gridregr, param_grid = grid_search, cv = 5)
gridcv_GDT = gridcv_GDT.fit(X_train,y_train)
view_best_params_GDT = gridcv_GDT.best_params_

#Apply model on test data
predictions_GDT = gridcv_GDT.predict(X_test)

#R^2
GDT_r2 = r2_score(y_test, predictions_GDT)
#Calculate MSE
GDT_mse = mean_squared_error(y_test, predictions_GDT)
#Calculate MAPE
GDT_mape = MAPE(y_test, predictions_GDT)

print('Grid Search CV Decision Regressor Model Performance:')
print('Best Parameters = ',view_best_params_GDT)
print('R-squared = {:.2}'.format(GDT_r2))
print('MSE = ',(GDT_mse))
print('*****')

Grid Search CV Decision Regressor Model Performance:
Best Parameters = {'max_depth': 5}
R-squared = 0.21.
MSE = 0.3294611813627842
*****
```

Figure 5.16: Grid search CV Decision tree model with a depth of 5, having  $R^2$  value 0.21 and MSE 0.32.

### Random forest

```
In [514]: from sklearn.ensemble import RandomForestRegressor

RF_reg = RandomForestRegressor(n_estimators = 1000, random_state=0).fit(X_train,y_train)
RF_reg.get_params()

#Apply model on test data
predictions_RF = RF_reg.predict(X_test)

#R^2
RF_r2 = r2_score(y_test, predictions_RF)
#Calculating MSE
RF_mse = np.mean(( y_test - predictions_RF)**2)

print('Random Forest Regressor Model Performance:')
print('Default Parameters = ',RF_reg.get_params())
print('R-squared = {:.4}'.format(RF_r2))
print('MSE = ',RF_mse)
print('*****')

Random Forest Regressor Model Performance:
Default Parameters = {'bootstrap': True, 'criterion': 'mse', 'max_depth': None, 'max_features': 'auto',
x_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_
amples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 1000, 'n_jobs': 1, 'oob_score': False,
andom_state': 0, 'verbose': 0, 'warm_start': False}
R-squared = 0.2736.
MSE = 0.30330721444799214
*****
```

Figure 5.17: Default Random forest model with  $R^2$  value 0.27 and MSE 0.3

### Random forest random Search CV

```
In [515]: ##Random Search CV
from sklearn.model_selection import RandomizedSearchCV

RRF = RandomForestRegressor(random_state = 0)
n_estimator = list(range(1,20,2))
depth = list(range(1,100,2))

# Create the random grid
rand_grid = {'n_estimators': n_estimator,
             'max_depth': depth}

randomcv_rf = RandomizedSearchCV(RRF, param_distributions = rand_grid, n_iter = 5, cv = 5, random_s
randomcv_rf = randomcv_rf.fit(X_train,y_train)
predictions_RRF = randomcv_rf.predict(X_test)

view_best_params_RRF = randomcv_rf.best_params_

best_model = randomcv_rf.best_estimator_

predictions_RRF = best_model.predict(X_test)

#R^2
RRF_r2 = r2_score(y_test, predictions_RRF)
#Calculating MSE
RRF_mse = np.mean(( y_test - predictions_RRF)**2)

print('Random Search CV Random Forest Regressor Model Performance:')
print('Best Parameters = ',view_best_params_RRF)
print('R-squared = {:.2f}'.format(RRF_r2))
print('MSE = ',RRF_mse)

Random Search CV Random Forest Regressor Model Performance:
Best Parameters = {'n_estimators': 15, 'max_depth': 9}
R-squared = 0.26.
MSE = 0.30971718768557643
```

Figure 5.18: Random search CV Random forest model with a depth of 9 and n estimators 15. Having  $R^2$  value 0.26 and MSE 0.3.

### Random forest Grid Search CV

```
In [516]: ## Grid Search CV
from sklearn.model_selection import GridSearchCV

regr = RandomForestRegressor(random_state = 0)
n_estimator = list(range(11,20,1))
depth = list(range(5,15,2))

# Create the grid
grid_search = {'n_estimators': n_estimator,
               'max_depth': depth}

## Grid Search Cross-Validation with 5 fold CV
gridcv_rf = GridSearchCV(regr, param_grid = grid_search, cv = 5)
gridcv_rf = gridcv_rf.fit(X_train,y_train)
view_best_params_GRF = gridcv_rf.best_params_

#Apply model on test data
predictions_GRF = gridcv_rf.predict(X_test)

#R^2
GRF_r2 = r2_score(y_test, predictions_GRF)
#Calculating MSE
GRF_mse = np.mean(( y_test - predictions_GRF)**2)

print('Grid Search CV Random Forest Regressor Model Performance:')
print('Best Parameters = ',view_best_params_GRF)
print('R-squared = {:.0.2}'.format(GRF_r2))
print('MSE = ',(GRF_mse))

Grid Search CV Random Forest Regressor Model Performance:
Best Parameters = {'max_depth': 7, 'n_estimators': 16}
R-squared = 0.27.
MSE = 0.3027266645879698
```

Figure 5.19: Grid search CV Random forest model with a depth of 7 n estimators 16. Having  $R^2$  value 0.27 and MSE 0.3.

### Gradient Boost

```
In [517]: #Gradient Boost
from sklearn.ensemble import GradientBoostingRegressor

gbt = GradientBoostingRegressor(random_state= 0).fit(X_train,y_train)

predictions_gbt = gbt.predict(X_test)

gbt.get_params()

#R^2
GBR_r2 = r2_score(y_test, predictions_gbt)
#Calculate MSE
GBR_mse = mean_squared_error(y_test, predictions_gbt)

print('Gradient Boosting Regressor Model Performance:')
print('Default Parameters = ',gbt.get_params())
print('R-squared = {:.2f}'.format(GBR_r2))
print('MSE = ',GBR_mse)
print('*****')
```

Gradient Boosting Regressor Model Performance:  
Default Parameters = {'alpha': 0.9, 'criterion': 'friedman\_mse', 'init': None, 'learning\_rate': 0.1, 'loss': 'ls', 'max\_depth': 3, 'max\_features': None, 'max\_leaf\_nodes': None, 'min\_impurity\_decrease': 0.0, 'min\_impurity\_split': None, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'min\_weight\_fraction\_leaf': 0.0, 'n\_estimators': 100, 'presort': 'auto', 'random\_state': 0, 'subsample': 1.0, 'verbose': 0, 'warm\_start': False}  
R-squared = 0.26.  
MSE = 0.3106926099088413  
\*\*\*\*\*

Figure 5.20: Default Gradient Boosting model with max depth 3 and n estimator 100, has a  $R^2$  value 0.26 and MSE 0.3

### Random Search CV Gradient boosting

```
In [518]: ##Random Search CV
rGBR = GradientBoostingRegressor(random_state = 0)
#loss = ['ls','lad','huber','quantile']
n_estimator = list(range(50,150,10))
#max_feat = ['auto','sqrt','log2']
depth = list(range(1,10,2))

# Create the random grid
rand_GBT = {'loss': loss,
            'n_estimators': n_estimator,
            'max_features': max_feat,
            'max_depth': depth}

randomcv_gbt = RandomizedSearchCV(rGBR, param_distributions = rand_GBT, n_iter = 5, cv = 5, random_state=0)
randomcv_gbt = randomcv_gbt.fit(X_train,y_train)
predictions_GBT = randomcv_gbt.predict(X_test)

view_best_params_GBT = randomcv_gbt.best_params_

#R^2
rGBR_r2 = r2_score(y_test, predictions_GBT)
#Calculate MSE
rGBR_mse = mean_squared_error(y_test, predictions_GBT)

print('Random Search CV Gradient Boosting Regressor Model Performance:')
print('Best Parameters = ',view_best_params_GBT)
print('R-squared = {:.0.2}'.format(rGBR_r2))
print('MSE = ',rGBR_mse)
print('*****')

Random Search CV Gradient Boosting Regressor Model Performance:
Best Parameters = {'n_estimators': 60, 'max_depth': 3}
R-squared = 0.28.
MSE = 0.3021934297233443
*****
```

Figure 5.21: Random search CV Gradient Boosting model with a depth of 3 and n estimators 60. Having  $R^2$  value 0.28 and MSE 0.30.

## Gradient Boosting Grid Search CV

```
In [519]: ## Grid Search CV

gGBR = GradientBoostingRegressor(random_state=0)
#loss = ['ls', 'lad', 'huber', 'quantile']
n_estimator = list(range(40,80,5))
#max_feat = ['auto', 'sqrt', 'log2']
depth = list(range(1,5,1))

# Create the random grid
grid_GBT = {'loss': loss,
            'n_estimators': n_estimator,
            'max_features': max_feat,
            'max_depth': depth}

## Grid Search Cross-Validation with 5 fold CV
gridcv_GBT = GridSearchCV(gGBR, param_grid = grid_GBT, cv = 5)
gridcv_GBT = gridcv_GBT.fit(X_train,y_train)
view_best_params_gridGRF = gridcv_GBT.best_params_

#Apply model on test data
predictions_gridGBT = gridcv_GBT.predict(X_test)

#R^2
gGBR_r2 = r2_score(y_test, predictions_gridGBT)
#Calculate MSE
gGBR_mse = mean_squared_error(y_test, predictions_gridGBT)

#Calculate MAPE
gGBR_mape = MAPE(y_test, predictions_gridGBT)

print('Grid Search CV Gradient Boosting Regressor Model Performance:')
print('Best Parameters = ',view_best_params_gridGRF)
print('R-squared = {:.2}'.format(gGBR_r2))
print('MSE = ', gGBR_mse)
print('*****')

Grid Search CV Gradient Boosting Regressor Model Performance:
Best Parameters = {'max_depth': 2, 'n_estimators': 65}
R-squared = 0.31.
MSE = 0.28671413364552173
*****
```

Figure 5.22: Grid search CV Gradient Boosting model with a depth of 4 n estimators 45. Having  $R^2$  value 0.31 and MSE 0.28.

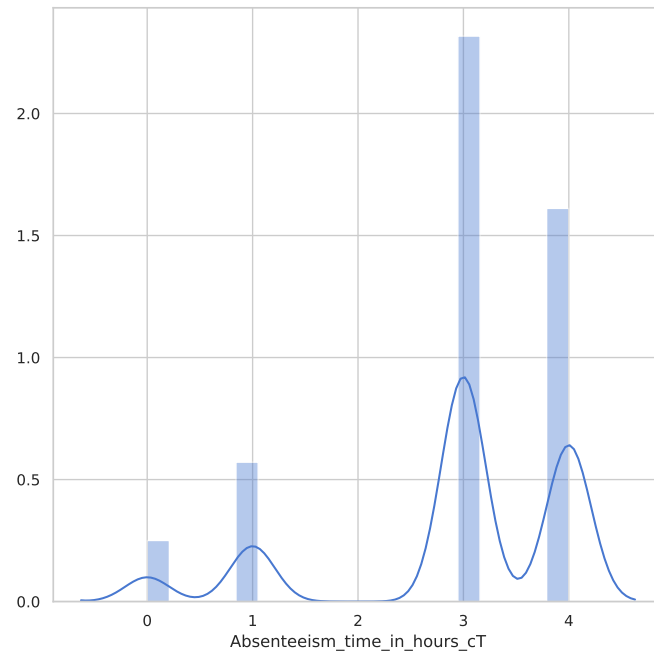


Figure 5.23: The distribution of the new binned target variable.

```
In [525]: from scipy.stats import chi2_contingency
#loop for chi square values
for i in obj_dtype:
    print(i)
    chi2, p, dof, ex = chi2_contingency(pd.crosstab(df['Absenteeism_time_in_hours'], df[i]))
    print(p)
```

ID  
7.705346449864873e-10  
Reason\_for\_absence  
3.038684341153542e-49  
Month\_of\_absence  
5.198801170459447e-09  
Disciplinary\_failure  
1.7506819169177113e-117  
Son  
2.226887327859972e-08  
Social\_drinker  
0.011539968647928535  
Pet  
0.12608552548493424

Figure 5.24: Chi squared test.

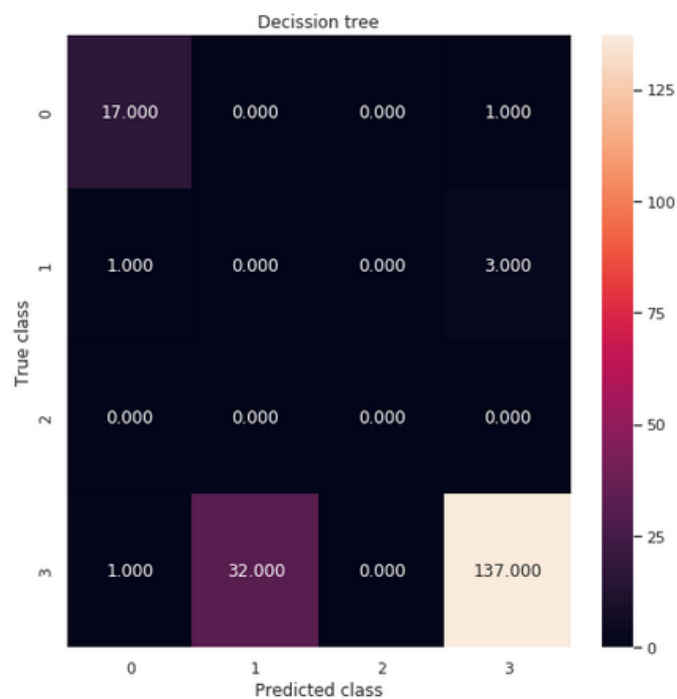


### Decision Tree Classifier

```
n [542]: from sklearn.tree import DecisionTreeClassifier
        clf = DecisionTreeClassifier(random_state=0).fit(X_train,y_train)

n [543]: y_pred = clf.predict(X_test)

n [553]: from sklearn.metrics import confusion_matrix
        from sklearn import metrics
        import seaborn as sns
        forest_cm = metrics.confusion_matrix(y_pred, y_test,[0,1,2,3])
        sns.heatmap(forest_cm, annot=True, fmt='.3f' )
        plt.ylabel('True class')
        plt.xlabel('Predicted class')
        plt.title('Decission tree')
        plt.savefig('dt.pdf')
```



```
n [545]: print('Accuracy: {:.3f}'.format(accuracy_score(y_test, clf.predict(X_test))))
```

Accuracy: 0.597

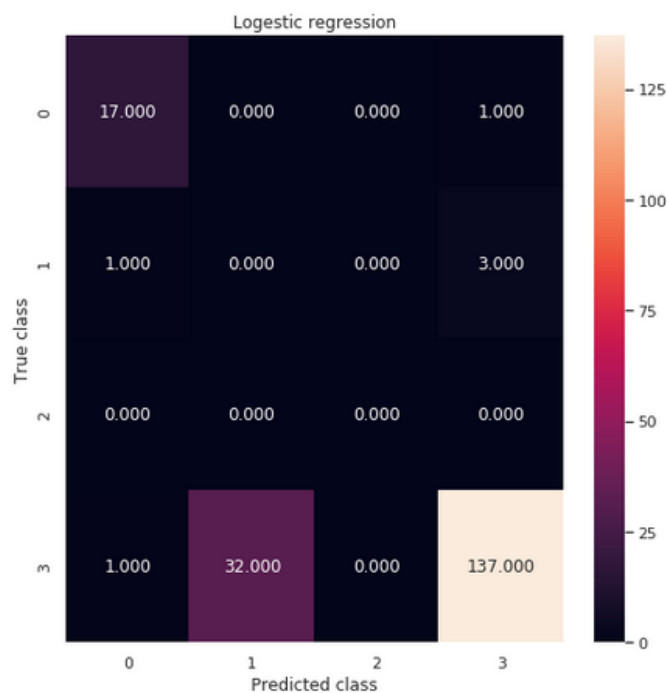
Figure 5.25: Decision tree multiclass classifier with an accuracy of 59.7%.

### Logistic Regression Multinomial classifier

```
In [554]: from sklearn.linear_model import LogisticRegression

clf = LogisticRegression(random_state=0, solver='lbfgs', multi_class='multinomial', max_iter=1000)
clf.fit(X_train, y_train)
y_pred=clf.predict(X_test)

from sklearn.metrics import confusion_matrix
from sklearn import metrics
import seaborn as sns
logr_cm = metrics.confusion_matrix(y_pred, y_test, [0,1,2,3])
sns.heatmap(logr_cm, annot=True, fmt='.3f')
plt.ylabel('True class')
plt.xlabel('Predicted class')
plt.title('Logestic regression')
plt.savefig('Logestic_regression')
```



```
In [555]: print('Accuracy: {:.3f}'.format(accuracy_score(y_test, clf.predict(X_test))))
```

Accuracy: 0.616

Figure 5.26: Logistic regression multi class classier with an curacy of 61.5%.

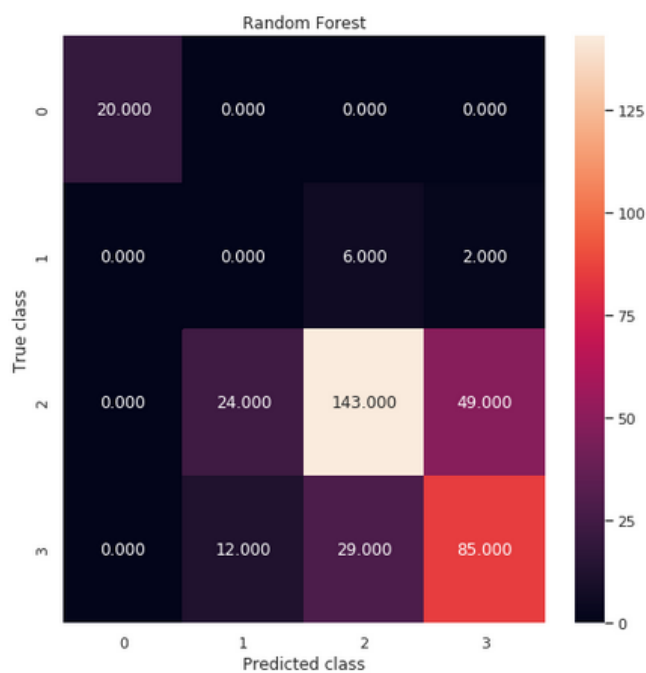
### Random Forest Multiclass classifier

```
In [556]: #Random Forest
from sklearn.ensemble import RandomForestClassifier

RF_model = RandomForestClassifier(n_estimators = 500).fit(X_train, y_train)
```

```
In [557]: RF_Predictions = RF_model.predict(X_test)
```

```
In [558]: from sklearn.metrics import confusion_matrix
from sklearn import metrics
import seaborn as sns
rf_cm = metrics.confusion_matrix(y_pred, RF_Predictions)
sns.heatmap(rf_cm, annot=True, fmt='.3f')
plt.ylabel('True class')
plt.xlabel('Predicted class')
plt.title('Random Forest')
plt.savefig('random_forest')
```



```
In [559]: print('Accuracy: {:.3f}'.format(accuracy_score(y_test, RF_model.predict(X_test))))
```

Accuracy: 0.670

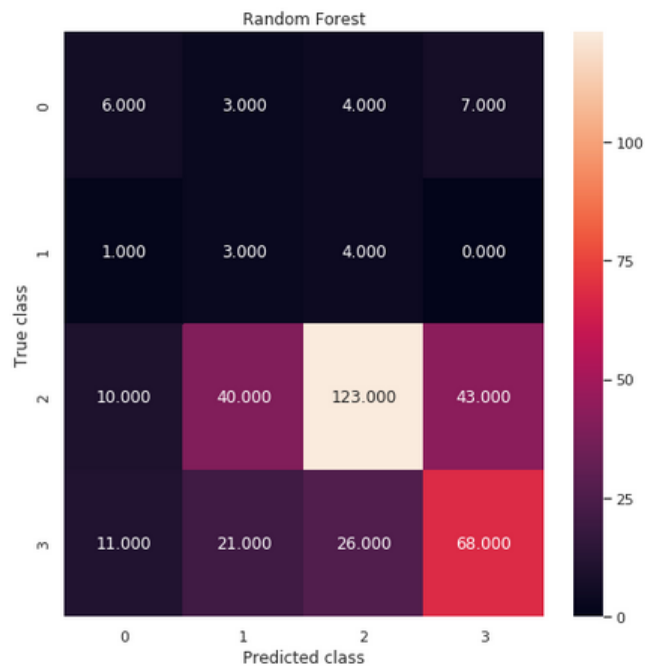
Figure 5.27: Random forest multiclass classifier with an accuracy of 67%.

## KNN

```
In [560]: from sklearn.neighbors import KNeighborsClassifier  
KNN_model = KNeighborsClassifier(n_neighbors = 3).fit(X_train, y_train)
```

```
In [561]: #predict test cases  
KNN_Predictions = KNN_model.predict(X_test)
```

```
In [562]: from sklearn.metrics import confusion_matrix  
from sklearn import metrics  
import seaborn as sns  
rf_cm = metrics.confusion_matrix(y_pred, KNN_Predictions)  
sns.heatmap(rf_cm, annot=True, fmt='.3f')  
plt.ylabel('True class')  
plt.xlabel('Predicted class')  
plt.title('Random Forest')  
plt.savefig('random_forest')
```



```
In [563]: print('Accuracy: {:.3f}'.format(accuracy_score(y_test, KNN_model.predict(X_test))))  
Accuracy: 0.505
```

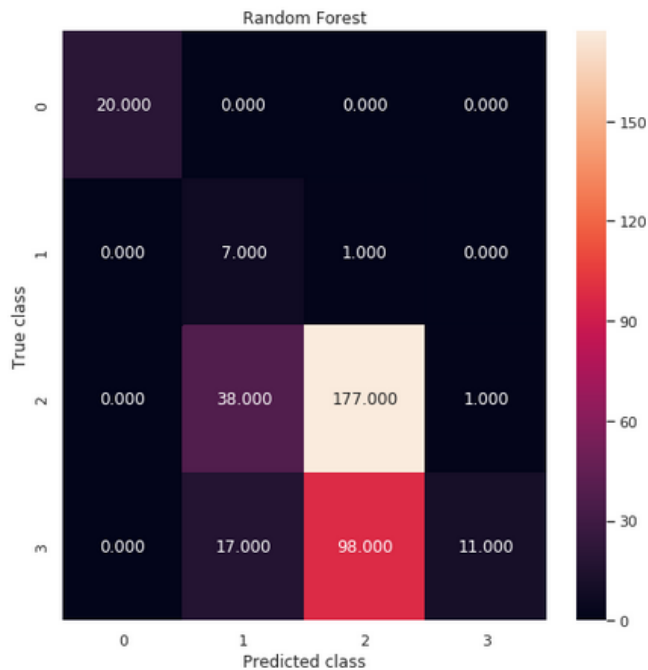
Figure 5.28: KNN classifier with an accuracy of 50.5%.

## Naive Bayes

```
In [564]: from sklearn.naive_bayes import GaussianNB
#Naive Bayes implementation
NB_model = GaussianNB().fit(X_train, y_train)
```

```
In [565]: #predict test cases
NB_Predictions = NB_model.predict(X_test)
```

```
In [566]: from sklearn.metrics import confusion_matrix
from sklearn import metrics
import seaborn as sns
rf_cm = metrics.confusion_matrix(y_pred, NB_Predictions)
sns.heatmap(rf_cm, annot=True, fmt='.3f' )
plt.ylabel('True class')
plt.xlabel('Predicted class')
plt.title('Random Forest')
plt.savefig('random_forest')
```



```
In [567]: print('Accuracy: {:.3f}'.format(accuracy_score(y_test, NB_model.predict(X_test))))
Accuracy: 0.473
```

Figure 5.29: Naive Bayes classifier with an curacy of 47.3%.

---

# Conclusion

---

## 6.1 Model Evaluation

Now that a few models have been created on our data set, it is required that a suitable evaluation matrix is selected to compare the different models. For the regression model, there are a few popular evaluation matrices which are commonly used. These matrices are: MAPE - Mean absolute percentage error, is the measure of accuracy as a percentage of error, MSE - Mean squared error, it is the mean squared errors and RMSE - Root mean squared error, it is the square root of the mean squared errors. For the classification model, we would be using accuracy.

MSE is selected as it gives us the goodness of fit of the model for the regression model. The smaller the MSE, the better the fit. R-squared values are also taken into account for our regression model as it explains as to how much of the variance of the target variable is explained. For the classification model, we would be using accuracy.

## 6.2 Model Selection

Table 6.1: Regression model performance after backward elimination

Model name	MSE	R-squared
Multivariate linear regression	0.247	0.41
Decision tree regressor	0.4	0.038
Random search cv Decision tree	0.329	0.21
Grid search cv Decision tree	0.341	0.18
Random forest	0.32	0.23
Random search cv Random forest	0.29	0.29
Grid search cv Random forest	0.29	0.29
Gradient boosting	0.3	0.28
Random search cv Gradient boosting	0.31	0.25
Grid search cv Gradient boosting	0.30	0.27

From table 6.1, 6.2 & 6.3 we can see that the best model with good performance is the multiclass random forest classifier model. In this model, the target variable is binned to form 4 separate categories. Due to this, we are able to get an accuracy of 67%. Hence, this model is selected.

Table 6.2: Regression model performance after backward elimination and PCA

Model name	MSE	R-squared
Multivariant linear regression	0.29	0.28
Decision tree regressor	0.37	0.10
Random search cv Decision tree	0.41	0.015
Grid search cv Decision tree	0.32	0.21
Random forest	0.3	0.27
Random search cv Random forest	0.3	0.26
Grid search cv Random forest	0.3	0.27
Gradient boosting	0.3	0.26
Random search cv Gradient boosting	0.3	0.28
Grid search cv Gradient boosting	0.28	0.31

Table 6.3: Classification model performance

Model name	Accuracy
Decision tree regressor	59.7%
Logestic regression	61.6%
Random Forest	67%
KNN	50.5%
Naive Naves	0.47

### 6.3 What changes company should bring to reduce the number of absenteeism?

From our analysis we were able to see that employees who had only an high school degree tend to be absent more often. It was also noted that people below the age of 30 and having service time less than 8 years tend to be absent often. Also when employees have medium work load i.e. between 250 - 300 average work load per day are often absent and people who have more or lesser work are not absent often. Also people with out children or pets often are absent compared to the rest. Therefor changes to be made is more college graduates to be hired in place of high school degree holders. People above the age of 30 need to be hired and the company need to retain people after they have 16 years of service. While hiring the employer should check to see if the employee or candidate has kids or pets and higher accordingly. Correct amount of work must be given to people. As it is noted people with work load 250-300 are absent often this could imply that the work load is less.

### 6.4 How much losses every month can we project in 2011 if same trend of absenteeism continues?

As we don't have the data for 2011, we shall use the given data to calculate the loss in 2011 assuming the trend remains same. We shall calculate the loss of work in time (hrs) which would be the total sum of absenteeism time in hours for each month respectively. We shall also calculate the loss in work load. Assuming work load average per day is the target workload for that day we shall

calculate its loss due to absenteeism time in hours by the formula given below

$$lossinwork = \frac{workload/day}{24} * Absenteeisminhours \quad (6.1)$$

Where 24 is the total number of hours in a day. The loss that will be incurred in 2011 is shown in figure 6.1

Total Absenteeism time in a month (hrs) Work loss per month		
Month_of_absence		
1.0	171.0	2252.833333
2.0	279.0	3164.875000
3.0	452.0	5281.916667
4.0	240.0	2728.875000
5.0	273.0	2794.166667
6.0	249.0	2815.083333
7.0	392.0	4129.750000
8.0	248.0	2438.000000
9.0	205.0	2298.666667
10.0	297.0	3326.708333
11.0	265.0	3136.833333
12.0	204.0	2208.291667

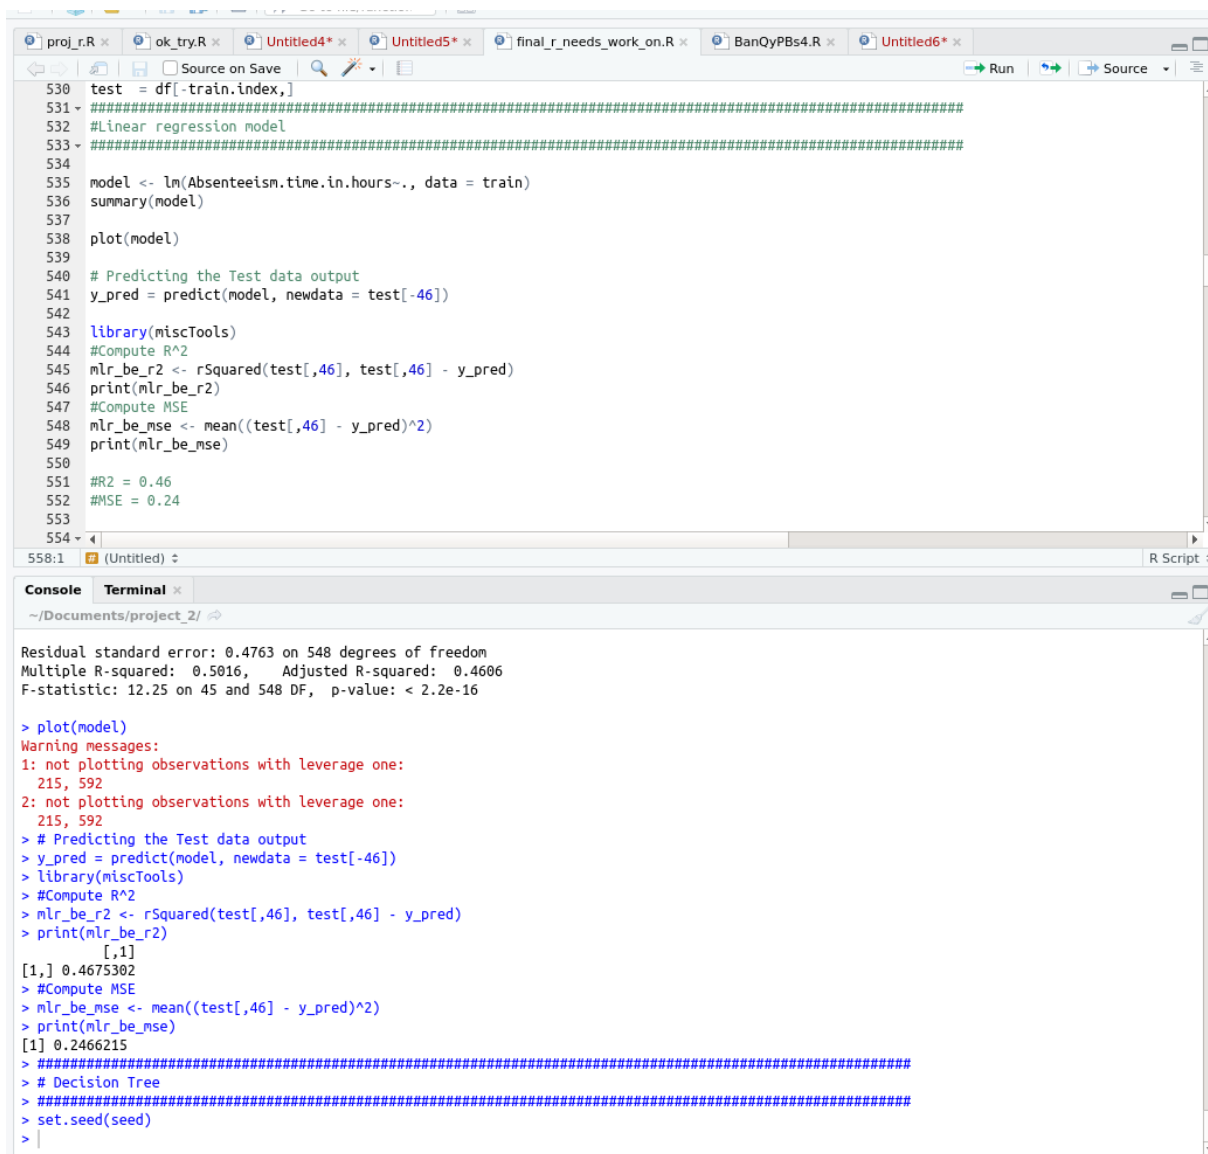
Figure 6.1: The loss that will be incurred in 2011



# **Appendices**

## Appendix A

# R code snapshots:



```
530 test = df[-train.index,]
531 #####
532 #Linear regression model
533 #####
534
535 model <- lm(Absenteeism.time.in.hours~., data = train)
536 summary(model)
537
538 plot(model)
539
540 # Predicting the Test data output
541 y_pred = predict(model, newdata = test[-46])
542
543 library(miscTools)
544 #Compute R^2
545 mlr_be_r2 <- rSquared(test[,46], test[,46] - y_pred)
546 print(mlr_be_r2)
547 #Compute MSE
548 mlr_be_mse <- mean((test[,46] - y_pred)^2)
549 print(mlr_be_mse)
550
551 #R2 = 0.46
552 #MSE = 0.24
553
554 <-
```

558:1 (Untitled) R Script

**Console** **Terminal**

~/Documents/project\_2/

Residual standard error: 0.4763 on 548 degrees of freedom  
Multiple R-squared: 0.5016, Adjusted R-squared: 0.4606  
F-statistic: 12.25 on 45 and 548 DF, p-value: < 2.2e-16

```
> plot(model)
Warning messages:
1: not plotting observations with leverage one:
   215, 592
2: not plotting observations with leverage one:
   215, 592
> # Predicting the Test data output
> y_pred = predict(model, newdata = test[-46])
> library(miscTools)
> #Compute R^2
> mlr_be_r2 <- rSquared(test[,46], test[,46] - y_pred)
> print(mlr_be_r2)
[1,] 0.4675302
> #Compute MSE
> mlr_be_mse <- mean((test[,46] - y_pred)^2)
> print(mlr_be_mse)
[1] 0.2466215
> #####
> # Decision Tree
> #####
> set.seed(seed)
>
```

Figure A.1: Linear regression model in r with output.

```

553
554 - #####
555 # Decision Tree
556 - #####
557 set.seed(seed)
558 #Building model
559 fit = rpart(Absenteeism.time.in.hours~., data = train, method = "anova")
560 #Variable importance
561 fit$variable.importance
562
563 # Predicting the Test data output
564 predictions_DT = predict(fit, test[-46])
565
566 #Compute R^2
567 dt_be_r2 <- rSquared(test[,46], test[,46] - predictions_DT)
568 print(dt_be_r2)
569 #Compute MSE
570 dt_be_mse <- mean((test[,46] - predictions_DT)^2)
571 print(dt_be_mse)
572
573 R2 = 0.51
574 #mse = 0.22
575 - #####
576 - ## Random Search #####
577 -
573:1 (Untitled)
R Script

```

```

~/Documents/project_2/
> #Building model
> fit = rpart(Absenteeism.time.in.hours~., data = train, method = "anova")
> #Variable importance
> fit$variable.importance
      Disciplinary.failure_1      Reason.for.absence_1      Transportation.expense      Reason.for.absence_2
      37.3801221             25.4563239             21.9437221             15.9377757
      Height                Age Distance.from.Residence.to.Work      ID_33
      14.1791039             11.7019322             9.6436484             5.3266177
      ID_10                  ID_5                ID_1             ID_13
      4.1659889              2.1961768             2.0220139             1.7358287
      ID_24                  Son_1                ID_23      Month.of.absence_5
      1.7051913              1.2435532             1.0414972             0.5865512
      Reason.for.absence_3      ID_26      Month.of.absence_8      ID_27
      0.5694331              0.4692410             0.3960150             0.2815652
      ID_6      Month.of.absence_12      ID_25
      0.2815652              0.2487106             0.1483980
> # Predicting the Test data output
> predictions_DT = predict(fit, test[-46])
> #Compute R^2
> dt_be_r2 <- rSquared(test[,46], test[,46] - predictions_DT)
> print(dt_be_r2)
      [,1]
[1,] 0.5137593
> #Compute MSE
> dt_be_mse <- mean((test[,46] - predictions_DT)^2)
> print(dt_be_mse)
[1] 0.2252097
>

```

Figure A.2: Decision tree model in r with output

The screenshot shows an RStudio interface with several open files. The active file contains R code for building and evaluating a Random Forest model. The console shows the execution of the code, including variable importance and model predictions.

```

645 #Random Forest
646 #####
647
648 #Building model
649 RF_model = randomForest(Absenteeism.time.in.hours ~ ., train, method = "anova", importance = TRUE)
650 #Prints out model information
651 print(RF_model)
652
653 #Predicting the Test data output
654 RF_Predictions = predict(RF_model, test[-46])
655
656 #Compute R^2
657 rf_be_r2 = rSquared(test[,46], test[,46] - RF_Predictions)
658 print(rf_be_r2)
659 #Compute MSE
660 rf_be_mse = mean((test[,46] - RF_Predictions)^2)
661 print(rf_be_mse)
662
663 #r2 = 0.50
664 #mse = 0.2
665 #####
666 ## Random Search #####
667 #####
668
669
573:1 (Untitled) R Script

```

The console output shows the following:

```

~/Documents/project_2/
> #Building model
> fit = rpart(Absenteeism.time.in.hours~., data = train, method = "anova")
> #Variable importance
> fit$variable.importance
      Disciplinary.failure_1      Reason.for.absence_1      Transportation.expense      Reason.for.absence_2
      37.3801221             25.4563239             21.9437221             15.9377757
      Height                Age Distance.from.Residence.to.Work             ID_33
      14.1791039             11.7019322             9.6436484             5.3266177
      ID_10                  ID_5                  ID_1                  ID_13
      4.1659889             2.1961768             2.0220139             1.7358287
      ID_24                  Son_1                  ID_23                  Month.of.absence_5
      1.7051913             1.2435532             1.0414972             0.5865512
      Reason.for.absence_3      ID_26                  Month.of.absence_8             ID_27
      0.5694331             0.4692410             0.3960150             0.2815652
      ID_6                   Month.of.absence_12      ID_25
      0.2815652             0.2487106             0.1483980
> # Predicting the Test data output
> predictions_DT = predict(fit, test[-46])
> #Compute R^2
> dt_be_r2 <- rSquared(test[,46], test[,46] - predictions_DT)
> print(dt_be_r2)
      [,1]
[1,] 0.5137593
> #Compute MSE
> dt_be_mse <- mean((test[,46] - predictions_DT)^2)
> print(dt_be_mse)
[1] 0.2252097
>

```

Figure A.3: Randeom forest model in r with output

```

1183 - #####
1184 #Randomforest
1185 - #####
1186 ## Building model
1187 model <- randomForest(Absenteeism.time.in.hours ~ ., data = train)
1188 # Predictions
1189 pred <- predict(model, newdata = test)
1190 #Confusion matrix
1191 cm=table(pred, test$Absenteeism.time.in.hours)
1192
1193 #Accuracy
1194 rf = sum(diag(cm))/nrow(test)
1195 print(rf)
1196 # accuracy 0.66
1197 - #####
1198 - ## Random Search #####
1199 - #####
1200
1201 # Create model with Random paramters 5 fold CV with 1 repeats
1202 control = trainControl(method="repeatedcv", number=5, repeats=3,search='random')
1203 set.seed(seed)
1204 rf_random = caret::train(Absenteeism.time.in.hours ~.,data = train, method="rf", metric= "Accuracy", tuneLength =10, trControl=control)
1205
1206 #print out summary of the model
1207
1203:1 (Untitled)
R Script

```

```

~/Documents/project 2/
Chi-squared approximation may be incorrect
8: In chisq.test(table(factor_data$Absenteeism.time.in.hours, factor_data[, :
Chi-squared approximation may be incorrect
9: In chisq.test(table(factor_data$Absenteeism.time.in.hours, factor_data[, :
Chi-squared approximation may be incorrect
> ## Dimension Reduction
> df = subset(df, select = -c(Body.mass.index, Day.of.the.week,Social.drinker))
> seed = 1234
> set.seed(seed)
> train.index = createDataPartition(df$Absenteeism.time.in.hours, p = .80, list = FALSE)
> train = df[ train.index,]
> test = df[-train.index,]
> #####
> #Randomforest
> #####
> ## Building model
> model <- randomForest(Absenteeism.time.in.hours ~ ., data = train)
> # Predictions
> pred <- predict(model, newdata = test)
> #Confusion matrix
> cm=table(pred, test$Absenteeism.time.in.hours)
> #Accuracy
> rf = sum(diag(cm))/nrow(test)
> print(rf)
[1] 0.6666667
> # Create model with Random paramters 5 fold CV with 1 repeats
> control = trainControl(method="repeatedcv", number=5, repeats=3,search='random')
>

```

Figure A.4: Randeom forest classifier in r

```

1228 #Multinomial logistic regression
1229 #####
1230
1231 Y = subset(df, select = c(Absenteeism.time.in.hours))
1232 X = subset(df, select = c(ID,Reason.for.absence,Month.of.absence,Transportation.expense, Distance.from.Residence.to.Work,
1233   Service.time ,Age,Work.load.Average.day.,Hit.target,Disciplinary.failure,Son,
1234   Social.smoker ,Pet ,Weight,Height))
1235
1236 # Function Creates the dummy variables and drops the first dummy variable
1237 X <- fastDummies::dummy_cols(X, remove_first_dummy = TRUE)
1238
1239 ## Dropping the original categorical variable
1240 #df = subset(df, select = -c(Pet,Son, Social.drinker,Social.smoker, Month.of.absence, Reason.for.absence, ID))
1241 X = subset(X, select = -c(Social.smoker,Pet,Son, Month.of.absence, Reason.for.absence, ID,Disciplinary.failure))
1242
1243 logit_df = data.frame(X, Y)
1244 #logit_df$out = relevel(logit_df$Absenteeism.time.in.hours,ref='1')
1245 seed = 1234
1246 set.seed(seed)
1247 train.index = createDataPartition(logit_df$Absenteeism.time.in.hours, p = .80, list = FALSE)
1248 logit_train = logit_df[train.index,]
1249 logit_test = logit_df[-train.index,]
1250
1251 mod <- multinom(Absenteeism.time.in.hours~., data = logit_train)
1252 summary(mod)
1253
1254 logit_predit =predict(mod,logit_test[1:70])
1255
1256 cm = table(logit_predit,logit_test$Absenteeism.time.in.hours)
1257 #Accuracy
1258 mlr = sum(diag(cm))/nrow(logit_test)
1259 print(mlr)
1260 # accuracy 0.66
1261 #####
1262
1260:1 (Untitled)

```

```

~/Documents/project_2/
3 0.3764703 0.4065824 1.1236916 0.2177173 0.2123974 1.422084e-01 2.404615e-01 0.2907838 0.2463259
1 0.2410002 0.2969574 3.126528e-01 3.050011e-11
2 0.2280600 0.2677925 6.340936e-23 1.992129e-01
3 0.2195214 0.3011260 3.126528e-01 1.992129e-01

Residual Deviance: 662.4649
AIC: 1016.465
Warning message:
In sqrt(diag(vcov)) : NaNs produced
> logit_predit =predict(mod,logit_test[1:70])
> cm = table(logit_predit,logit_test$Absenteeism.time.in.hours)
> #Accuracy
> mlr = sum(diag(cm))/nrow(logit_test)
> print(mlr)
[1] 0.6666667
>

```

Figure A.5: Logistic regression model in r

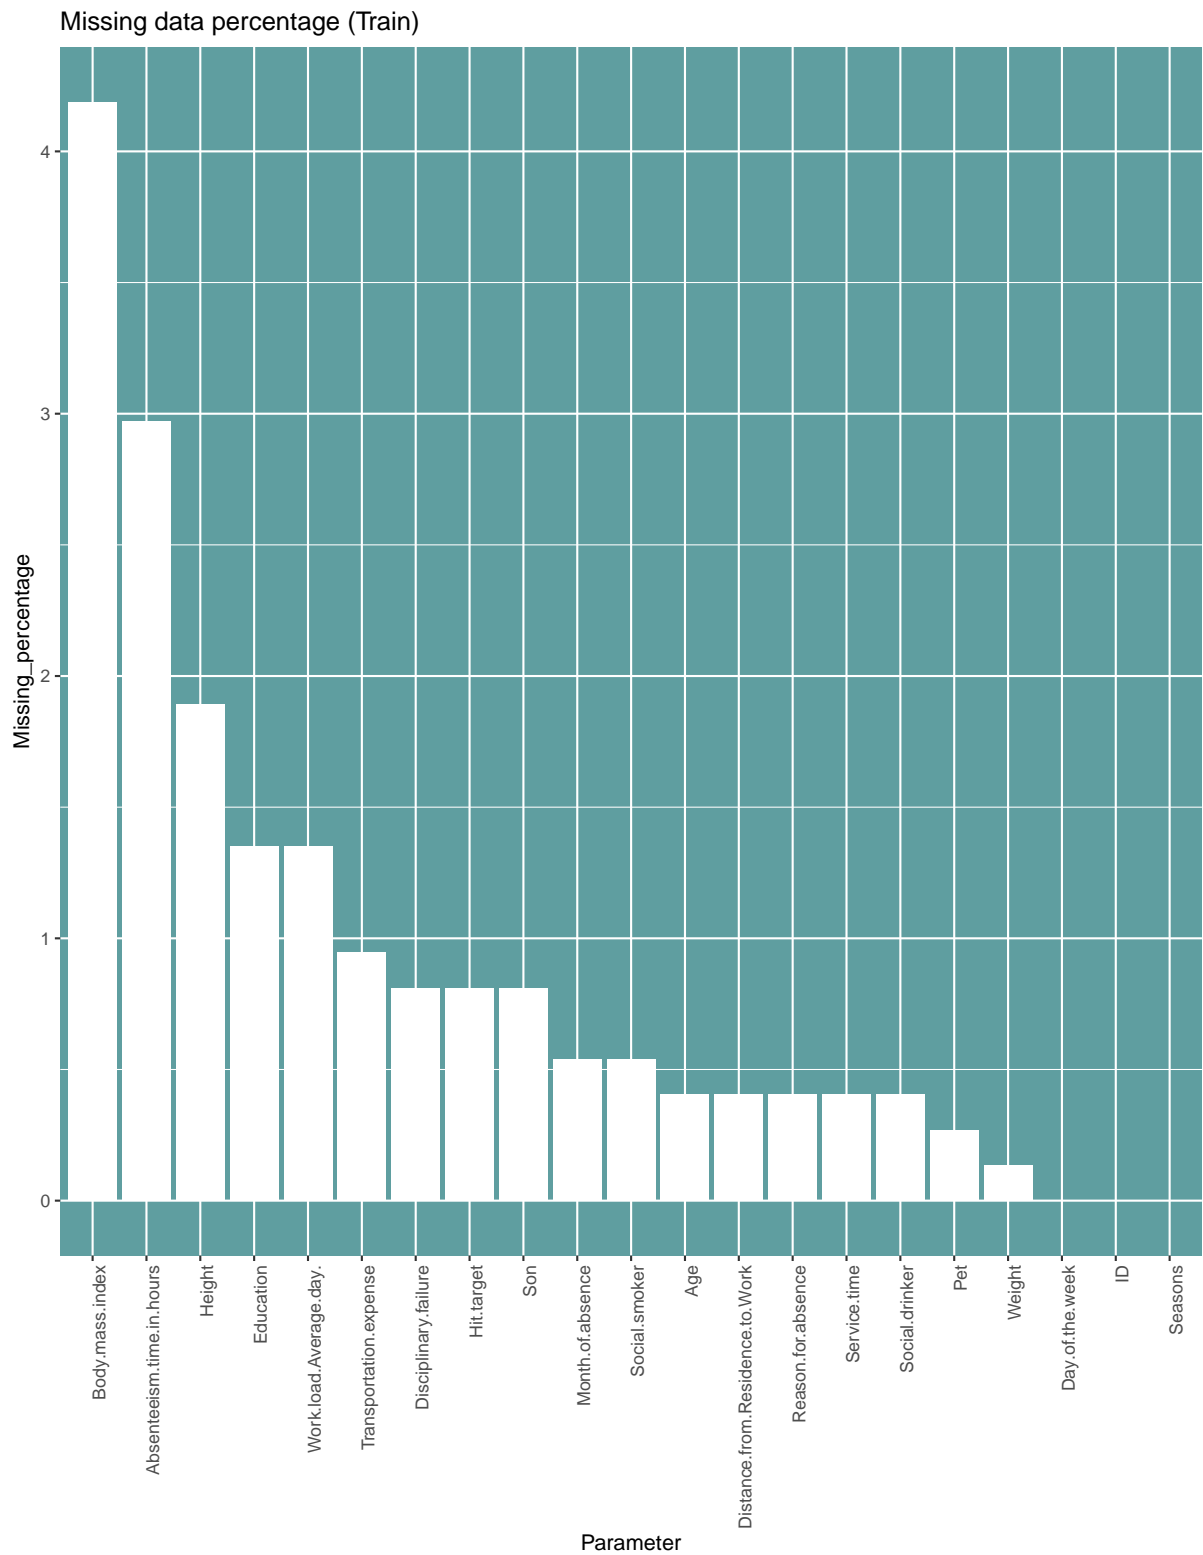


Figure A.6: missing value plot.